

Learning with Guarantees

Parametrising Neural Feedback Policies with Stability & Robustness Certificates



Nicholas H. Barbara

Doctor of Philosophy – Defence

20 October 2025



THE UNIVERSITY OF
SYDNEY

Australian Centre
for Robotics

Outline

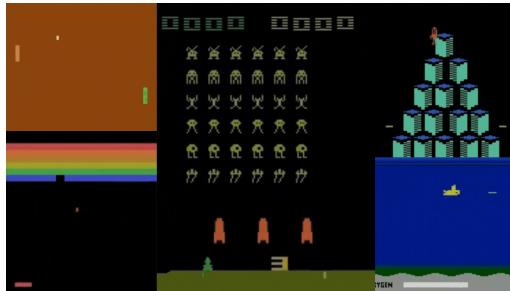
1. Motivation
2. Background: robust neural networks
3. Lipschitz-bounded policy networks
4. React to surprises with Youla-REN (Theory)
5. React to surprises with Youla-REN (Application)
6. Robust recurrent deep networks
7. Conclusions

Outline

1. Motivation
2. Background: robust neural networks
3. Lipschitz-bounded policy networks
4. React to surprises with Youla-REN (Theory)
5. React to surprises with Youla-REN (Application)
6. Robust recurrent deep networks
7. Conclusions

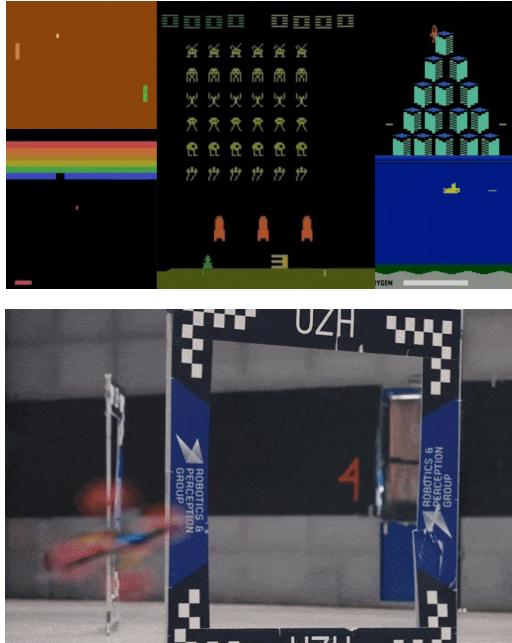
Deep Reinforcement Learning (RL)

Deep Reinforcement Learning (RL)



[1] Mnih et al., Human-level control through deep reinforcement learning. *Nature* (2015).

Deep Reinforcement Learning (RL)

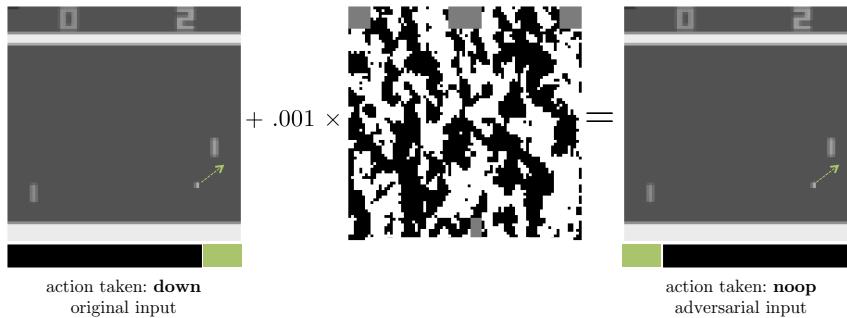


- [1] Mnih et al., Human-level control through deep reinforcement learning. *Nature* (2015).
- [2] Foehn et al., Agilicious: Open-Source and Open-Hardware Agile Quadrotor for Vision-Based Flight. *Science Robotics* (2022).
- [3] Miki et al., Learning robust perceptive locomotion for quadrupedal robots in the wild. *Science Robotics* (2022).
- [4] Siekmann et al., Sim-to-real learning of all common bipedal gaits via periodic reward composition. *ICRA* (2021).

No Closed-Loop Guarantees

No Closed-Loop Guarantees

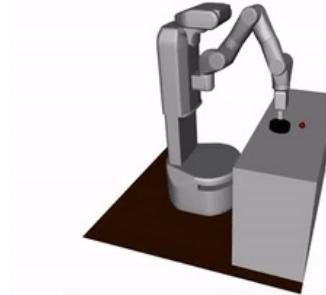
- Deep RL is not robust to small perturbations



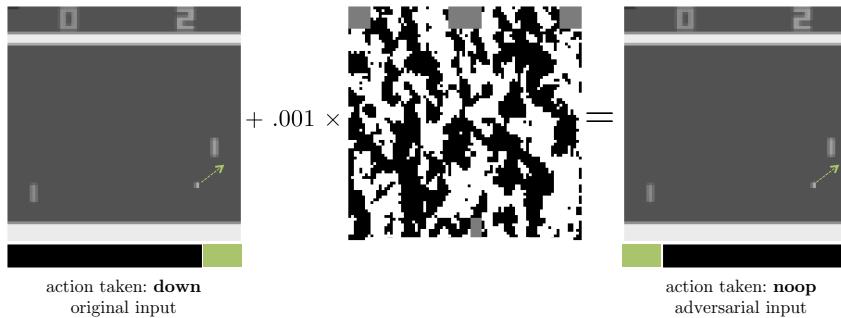
[5] Huang et al., Adversarial attacks on neural network policies. *ICLR* (2017).

No Closed-Loop Guarantees

- Deep RL is not robust to small perturbations
- Can improve robustness with empirical methods



<https://youtu.be/2AGQolc8CiQ?t=50>

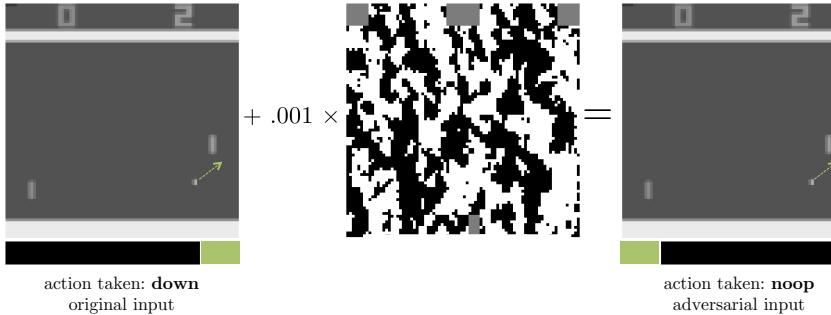


[5] Huang et al., Adversarial attacks on neural network policies. *ICLR* (2017).

[6] Peng et al., Sim-to-Real Transfer of Robotic Control with Dynamics Randomization. *ICRA* (2018).

No Closed-Loop Guarantees

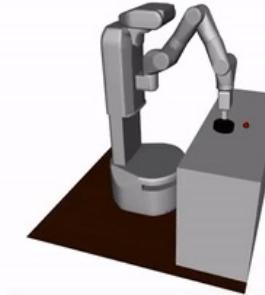
- Deep RL is not robust to small perturbations
- Can improve robustness with empirical methods
- ...but without *guarantees*, still closed-loop unstable!



[5] Huang et al., Adversarial attacks on neural network policies. *ICLR* (2017).

[6] Peng et al., Sim-to-Real Transfer of Robotic Control with Dynamics Randomization. *ICRA* (2018).

[7] Shi et al., Rethinking Robustness Assessment: Adversarial Attacks on Learning-based Quadrupedal Locomotion Controllers. *Robotics: Science and Systems* (2024).



<https://youtu.be/2AGQolc8CiQ?t=50>



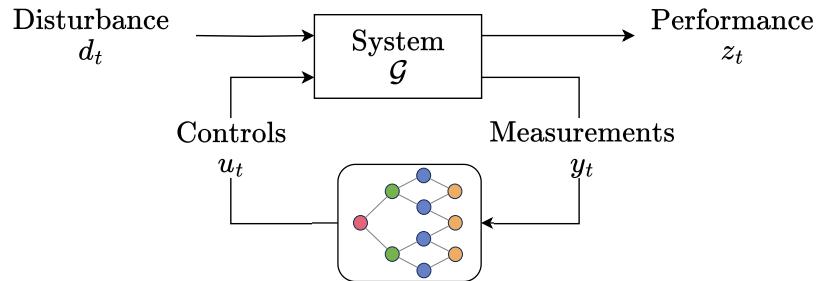
Black-Box Methods

Black-Box Methods

- Deep RL relies on black-box Neural Networks (NNs)

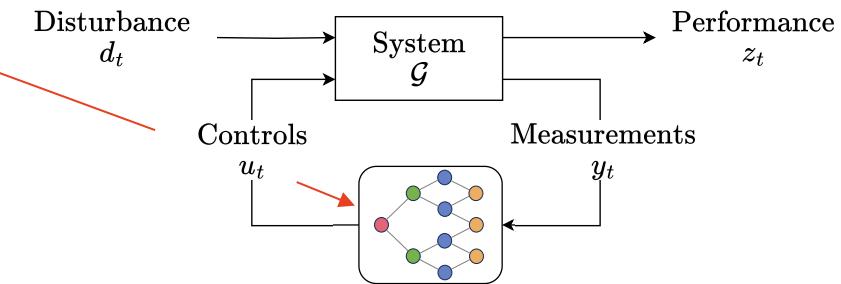
Black-Box Methods

- Deep RL relies on black-box Neural Networks (NNs)



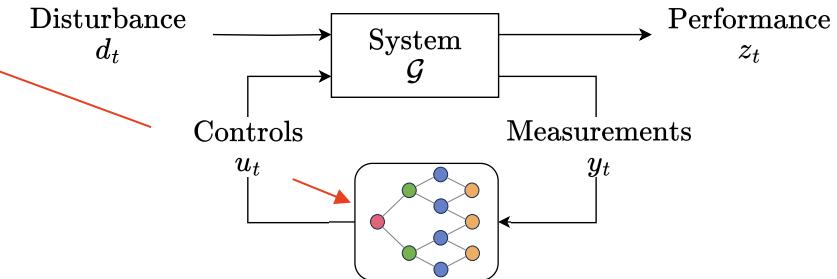
Black-Box Methods

- Deep RL relies on black-box Neural Networks (NNs)



Black-Box Methods

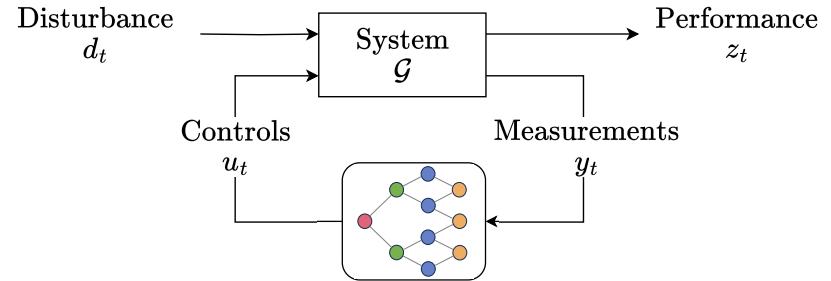
- Deep RL relies on black-box Neural Networks (NNs)
- Problem:
 - No prior system knowledge in the policy **architecture**
 - No prior system knowledge in the policy **parametrisation**
- ...so closed-loop stability is
not explicitly considered!



Question

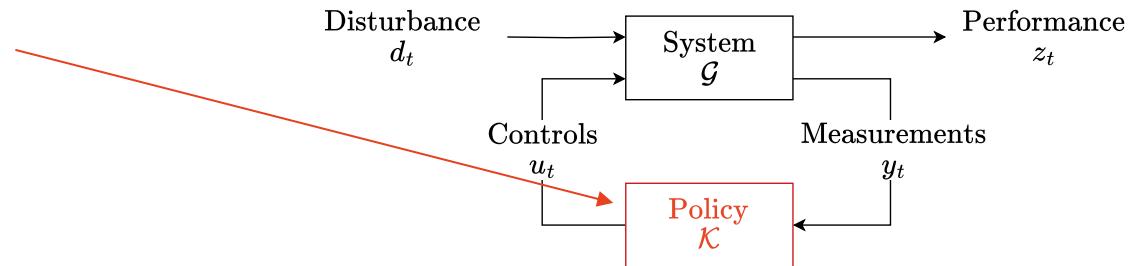
*Can we learn robustly-stabilising neural feedback policies
with **closed-loop guarantees** for nonlinear systems?*

Stable-by-Design Neural Feedback Policies



Stable-by-Design Neural Feedback Policies

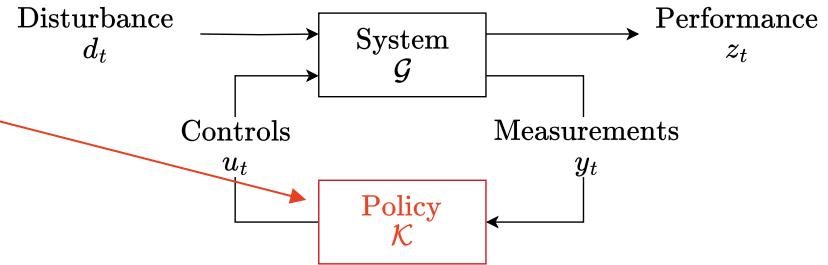
Construct a new *policy class*:



Stable-by-Design Neural Feedback Policies

Construct a new *policy class*:

1. Closed-loop robust stability by construction
2. Differentiable, scalable parametrisation
3. Expressive parametrisation

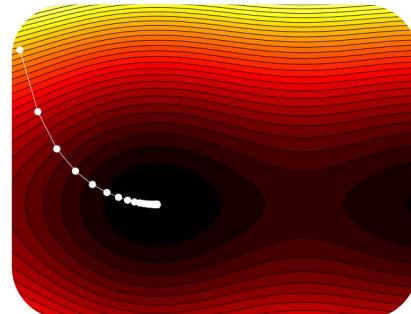
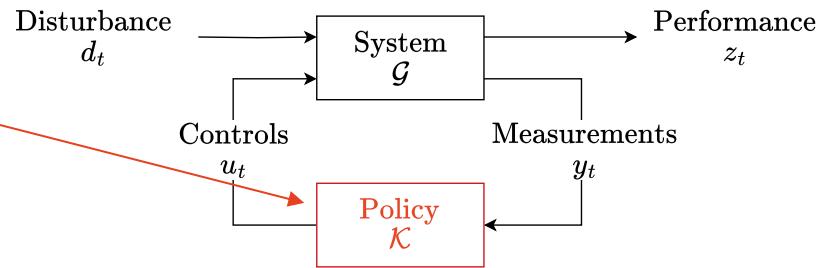


Stable-by-Design Neural Feedback Policies

Construct a new *policy class*:

1. Closed-loop robust stability by construction
2. Differentiable, scalable parametrisation
3. Expressive parametrisation

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools



Outline

1. Motivation
2. Background: robust neural networks
3. Lipschitz-bounded policy networks
4. React to surprises with Youla-REN (Theory)
5. React to surprises with Youla-REN (Application)
6. Robust recurrent deep networks
7. Conclusions

Robust Neural Networks

Robust Neural Networks

- Problems are caused by black-box NNs
- Sensitive to changes in their inputs, issues with internal instability

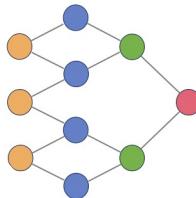
Robust Neural Networks

- Problems are caused by black-box NNs
- Sensitive to changes in their inputs, issues with internal instability
- Solution: use *robust neural networks*^[8]

[8] Manchester, Wang, & Barbara, Neural Networks in the Loop: Learning with Stability and Robustness Guarantees, accepted to *Ann. Rev. Control, Robotics, Autonomous Systems* (2025).

Robust Neural Networks

- Problems are caused by black-box NNs
- Sensitive to changes in their inputs, issues with internal instability
- Solution: use *robust neural networks*^[8]



Neural networks

[8] Manchester, Wang, & Barbara, Neural Networks in the Loop: Learning with Stability and Robustness Guarantees, accepted to *Ann. Rev. Control, Robotics, Autonomous Systems* (2025).

Robust Neural Networks

- Problems are caused by black-box NNs
- Sensitive to changes in their inputs, issues with internal instability
- Solution: use *robust neural networks*^[8]



Neural networks

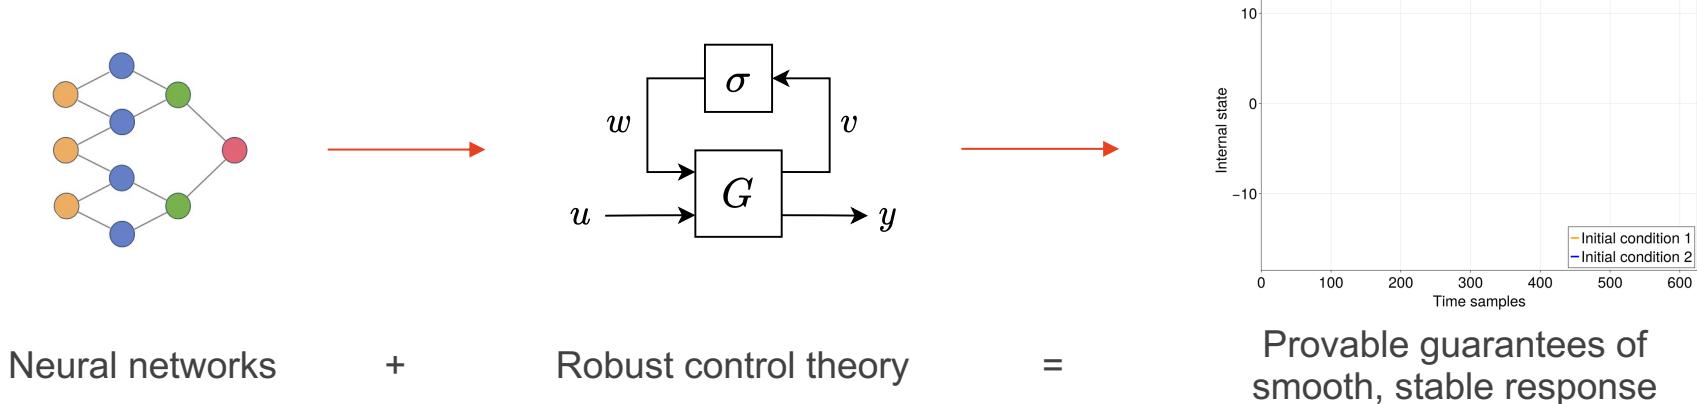
+

Robust control theory

[8] Manchester, Wang, & Barbara, Neural Networks in the Loop: Learning with Stability and Robustness Guarantees, accepted to Ann. Rev. Control, Robotics, Autonomous Systems (2025).

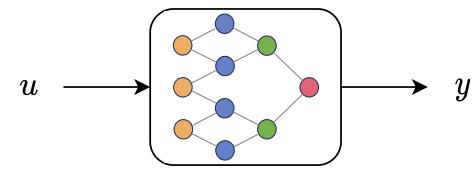
Robust Neural Networks

- Problems are caused by black-box NNs
- Sensitive to changes in their inputs, issues with internal instability
- Solution: use *robust neural networks*^[8]



[8] Manchester, Wang, & Barbara, Neural Networks in the Loop: Learning with Stability and Robustness Guarantees, accepted to Ann. Rev. Control, Robotics, Autonomous Systems (2025).

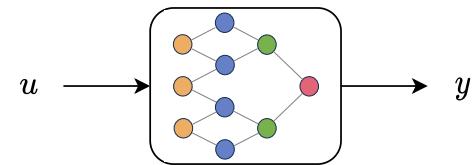
Input-Output Robustness: Lipschitz Continuity



Input-Output Robustness: Lipschitz Continuity

- Lipschitz systems respond smoothly to changes in their inputs:

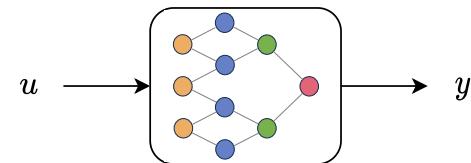
$$\|\Delta y\|_2 \leq \gamma \|\Delta u\|_2, \quad \gamma \geq 0$$



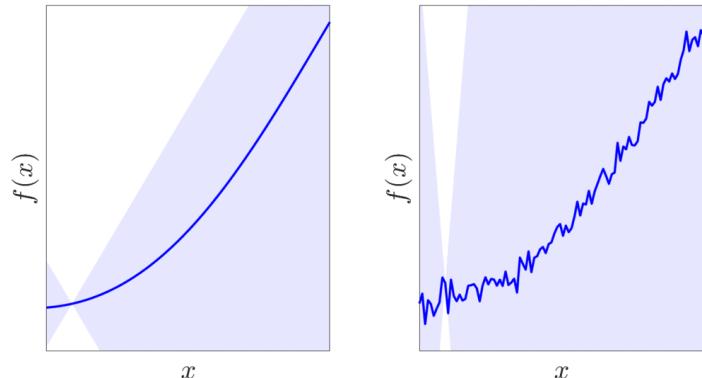
Input-Output Robustness: Lipschitz Continuity

- Lipschitz systems respond smoothly to changes in their inputs:

$$\|\Delta y\|_2 \leq \gamma \|\Delta u\|_2, \quad \gamma \geq 0$$



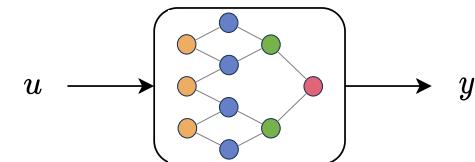
- Small Lipschitz constant \rightarrow smooth response
- Large Lipschitz constant \rightarrow sensitive to perturbations



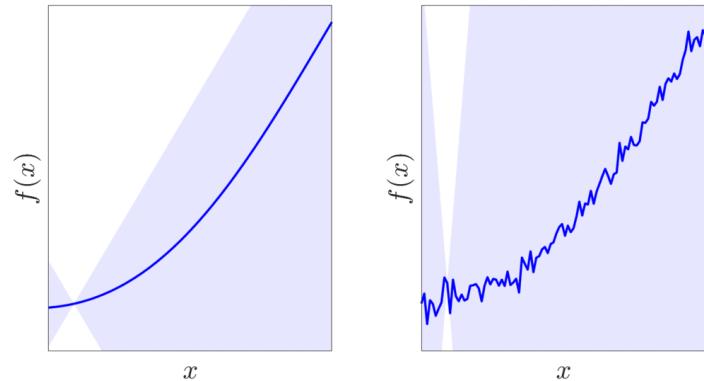
Input-Output Robustness: Lipschitz Continuity

- Lipschitz systems respond smoothly to changes in their inputs:

$$\|\Delta y\|_2 \leq \gamma \|\Delta u\|_2, \quad \gamma \geq 0$$



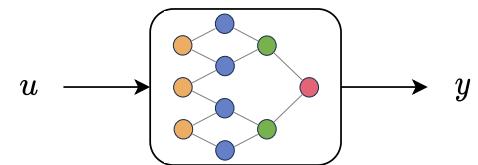
- Small Lipschitz constant \rightarrow smooth response
- Large Lipschitz constant \rightarrow sensitive to perturbations
- Generalise to other I/O properties with IQCs^[10]



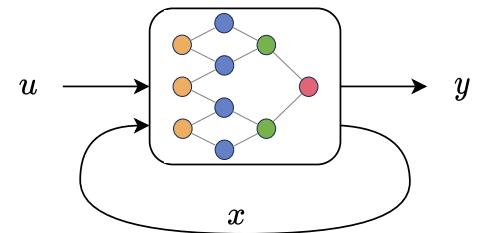
[9] Pauli et al., Neural Network Training Under Semidefinite Constraints, *CDC* (2022).

[10] Megretski & Rantzer, System Analysis via Integral Quadratic Constraints, *TAC* (1997).

Internal Stability: Contraction

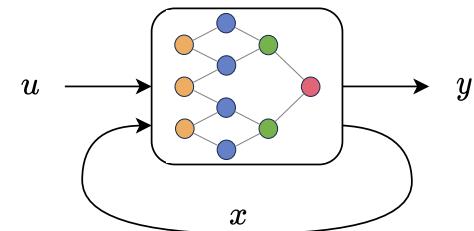


Internal Stability: Contraction



Internal Stability: Contraction

- Contracting systems forget their initial conditions exponentially

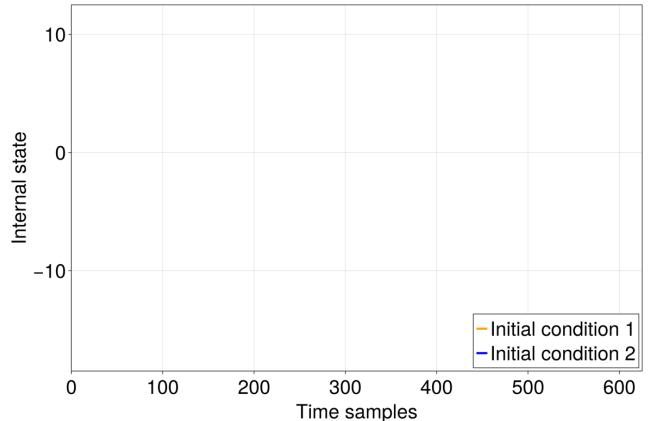
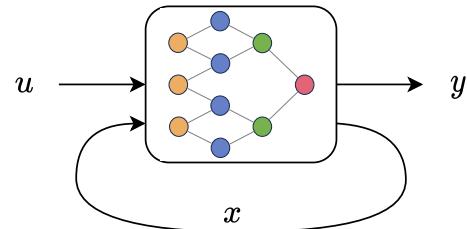


Internal Stability: Contraction

- Contracting systems forget their initial conditions exponentially
- For a discrete-time contracting system ($\Delta u = 0$)

$$|\Delta x_t| \leq \beta \alpha^t |\Delta x_0|, \quad \alpha \in [0, 1), \quad \beta \geq 0$$

- I.e., incremental exponential stability



[11] Lohmiller & Slotine, On Contraction Analysis for Nonlinear Systems. *Automatica* (1998).

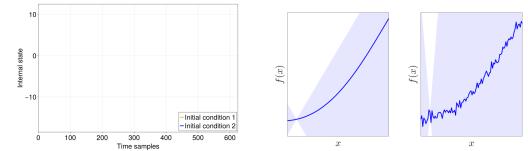
Stable & Robust Neural Networks

- Our suite of robust neural networks @ the ACFR^[8]:

[8] Manchester, Wang, & Barbara, Neural Networks in the Loop: Learning with Stability and Robustness Guarantees, *accepted to Ann. Rev. Control Robotics & Autonomous Systems*. (2025).

Stable & Robust Neural Networks

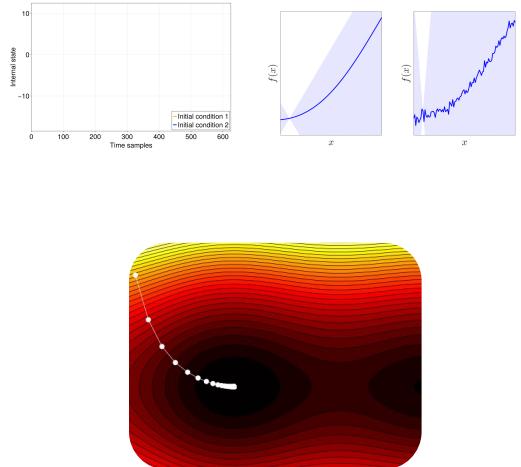
- Our suite of robust neural networks @ the ACFR^[8]:
 - **Static** (feedforward) and **dynamic** (recurrent) models
 - **Built-in guarantees** of incremental stability (contraction) & robustness (Lipschitz, IQC)



[8] Manchester, Wang, & Barbara, Neural Networks in the Loop: Learning with Stability and Robustness Guarantees, accepted to *Ann. Rev. Control Robotics & Autonomous Systems*. (2025).

Stable & Robust Neural Networks

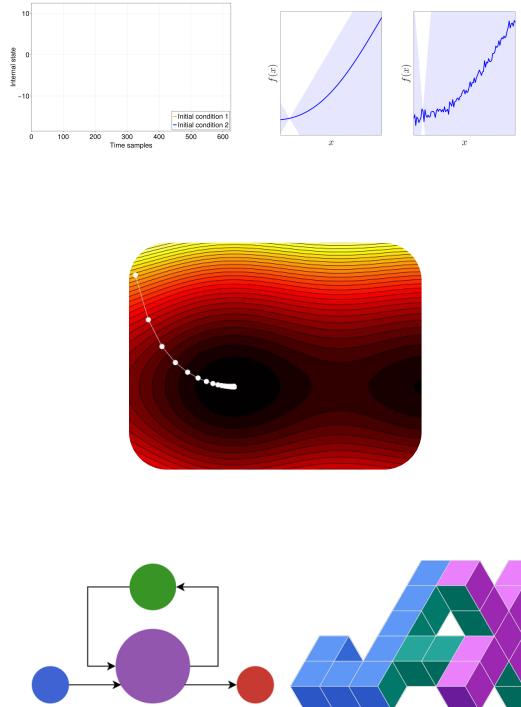
- Our suite of robust neural networks @ the ACFR^[8]:
 - **Static** (feedforward) and **dynamic** (recurrent) models
 - **Built-in guarantees** of incremental stability (contraction) & robustness (Lipschitz, IQC)
 - **Directly parametrised** to allow unconstrained optimisation
 - **Universal approximators** of all contracting & Lipschitz systems



[8] Manchester, Wang, & Barbara, Neural Networks in the Loop: Learning with Stability and Robustness Guarantees, accepted to *Ann. Rev. Control Robotics & Autonomous Systems*. (2025).

Stable & Robust Neural Networks

- Our suite of robust neural networks @ the ACFR^[8]:
 - **Static** (feedforward) and **dynamic** (recurrent) models
 - **Built-in guarantees** of incremental stability (contraction) & robustness (Lipschitz, IQC)
 - **Directly parametrised** to allow unconstrained optimisation
 - **Universal approximators** of all contracting & Lipschitz systems
- Open-source and ready for use in Julia and Python (JAX)^[8,12]



[8] Manchester, Wang, & Barbara, Neural Networks in the Loop: Learning with Stability and Robustness Guarantees, accepted to *Ann. Rev. Control Robotics & Autonomous Systems*. (2025).

[12] Barbara et al., RobustNeuralNetworks.jl: a Package for Machine Learning and Data-Driven Control with Certified Robustness, *JuliaCon Proceedings* (2025).

Examples of Robust Neural Networks

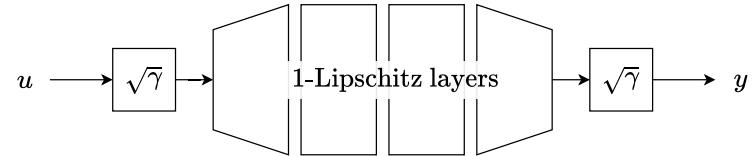
Two main robust NN architectures used in this talk:

Examples of Robust Neural Networks

Two main robust NN architectures used in this talk:

1. Sandwich network^[13]

- Static, feedforward DNNs (MLP, CNN, etc...)
- Certified **Lipschitz** bound



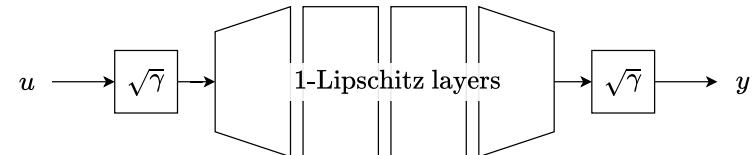
[13] Wang & Manchester, Direct Parameterization of Lipschitz-Bounded Deep Networks, *ICML* (2023).

Examples of Robust Neural Networks

Two main robust NN architectures used in this talk:

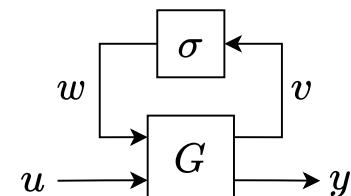
1. Sandwich network^[13]

- Static, feedforward DNNs (MLP, CNN, etc...)
- Certified **Lipschitz** bound



2. Recurrent Equilibrium Network (REN)^[14]

- Recurrent / dynamic models
- Certified **contraction** and **Lipschitz** bounds
- **Universal approximator** for contracting & Lipschitz systems



[13] Wang & Manchester, Direct Parameterization of Lipschitz-Bounded Deep Networks, *ICML* (2023).

[14] Revay, Wang, & Manchester, Recurrent Equilibrium Networks: Flexible Dynamic Models with Guaranteed Stability and Robustness, *TAC* (2023).

Examples of Robust Neural Networks

Two main robust NN architectures used in this talk:

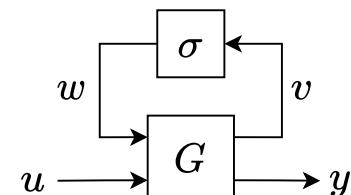
1. Sandwich network^[13]

- Static, feedforward DNNs (MLP, CNN, etc...)
- Certified **Lipschitz** bound



2. Recurrent Equilibrium Network (REN)^[14]

- Recurrent / dynamic models
- Certified **contraction** and **Lipschitz** bounds
- **Universal approximator** for contracting & Lipschitz systems



For more, see Manchester et al., *Ann. Rev. Control, Robotics, Autonomous Systems* (2025).

[13] Wang & Manchester, Direct Parameterization of Lipschitz-Bounded Deep Networks, *ICML* (2023).

[14] Revay, Wang, & Manchester, Recurrent Equilibrium Networks: Flexible Dynamic Models with Guaranteed Stability and Robustness, *TAC* (2023).

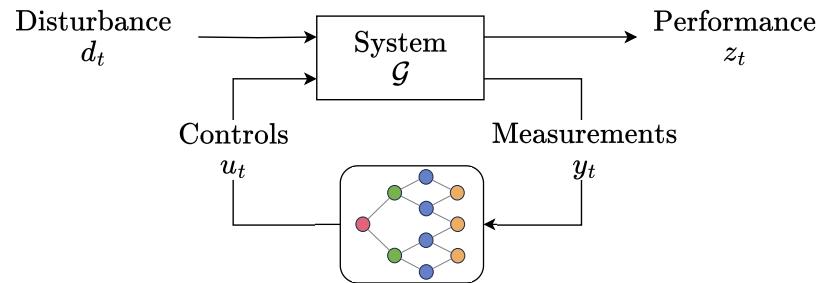
Outline

1. Motivation
2. Background: robust neural networks
3. Lipschitz-bounded policy networks
4. React to surprises with Youla-REN (Theory)
5. React to surprises with Youla-REN (Application)
6. Robust recurrent deep networks
7. Conclusions

Lipschitz-Bounded Policy Networks

Construct a new *policy class*:

1. Closed-loop robust stability by construction
2. Differentiable, scalable parametrisation
3. Expressive parametrisation



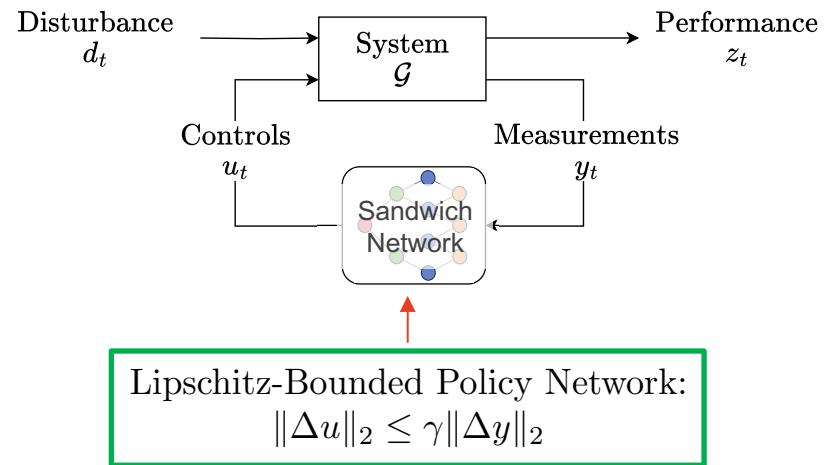
Directly learn controllers with robust stability guarantees using standard ML/deep RL tools

Lipschitz-Bounded Policy Networks

Construct a new *policy class*:

1. Closed-loop robust stability by construction
2. Differentiable, scalable parametrisation
3. Expressive parametrisation

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools

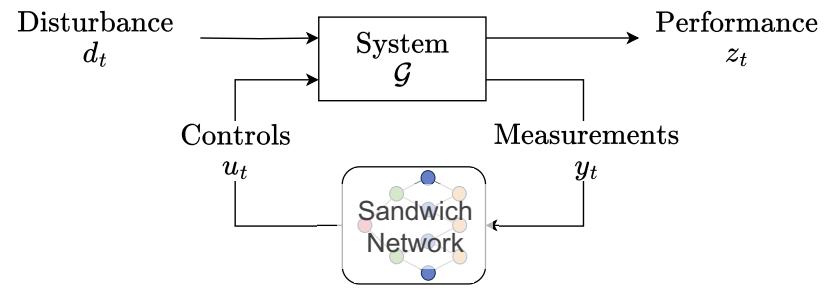


Lipschitz-Bounded Policy Networks

Construct a new *policy class*:

- X 1. Closed-loop robust stability by construction
- ✓ 2. Differentiable, scalable parametrisation
- ✓ 3. Expressive parametrisation

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools



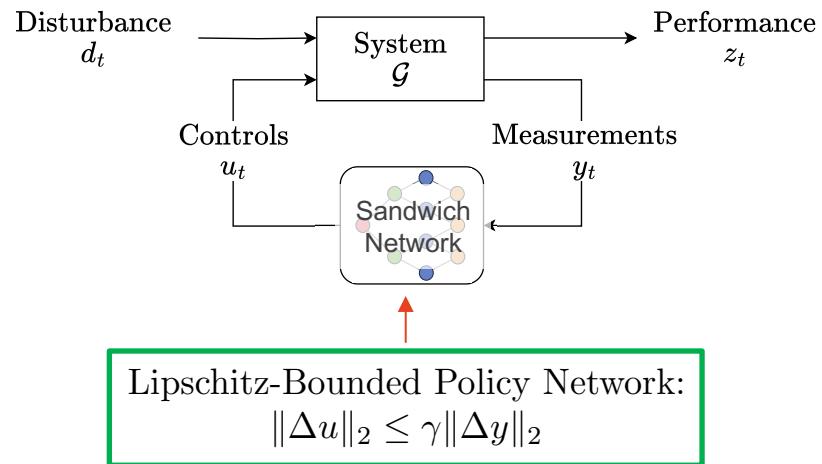
Lipschitz-Bounded Policy Network:
 $\|\Delta u\|_2 \leq \gamma \|\Delta y\|_2$

Lipschitz-Bounded Policy Networks

Construct a new *policy class*:

- 1. Robust to perturbations by construction
- 2. Differentiable, scalable parametrisation
- 3. Expressive parametrisation

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools

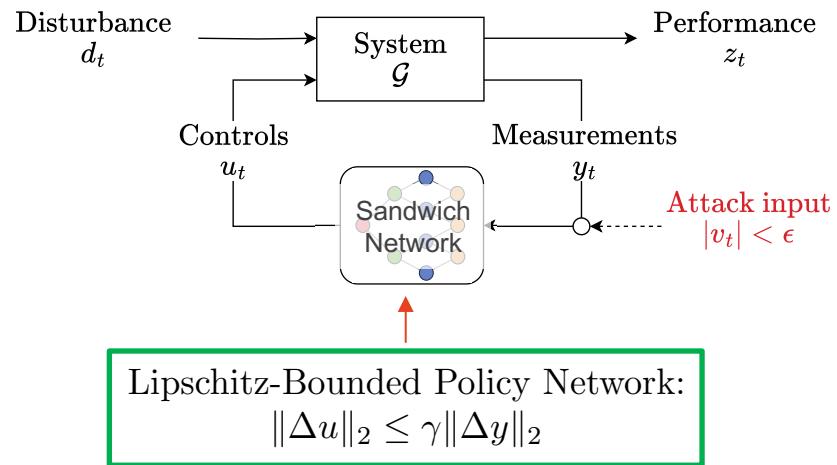


Lipschitz-Bounded Policy Networks

Construct a new *policy class*:

- 1. Robust to perturbations by construction
- 2. Differentiable, scalable parametrisation
- 3. Expressive parametrisation

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools

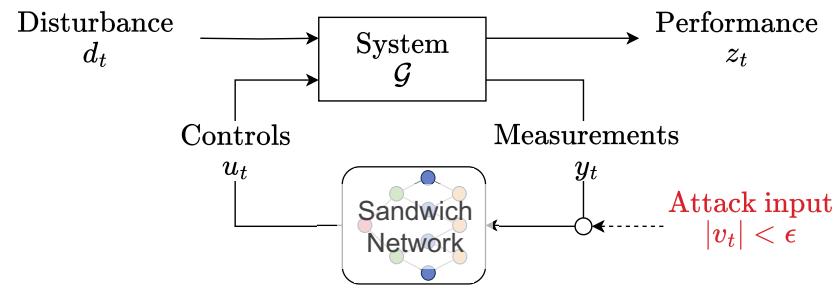


Lipschitz-Bounded Policy Networks

Construct a new *policy class*:

- ✓ 1. Robust to perturbations by construction
- ✓ 2. Differentiable, scalable parametrisation
- ✓ 3. Expressive parametrisation

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools



Lipschitz-Bounded Policy Network:
 $\|\Delta u\|_2 \leq \gamma \|\Delta y\|_2$

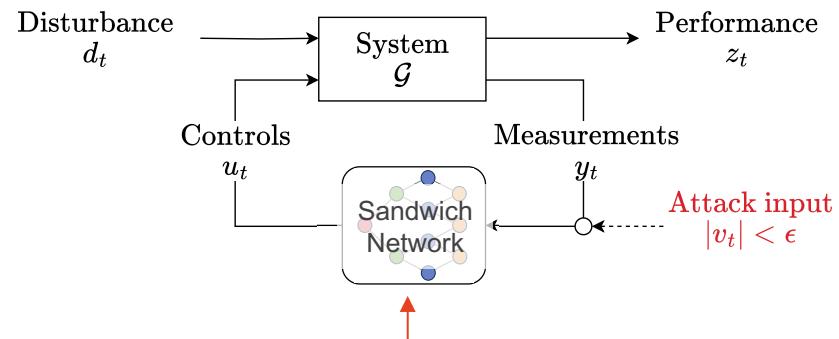
$$u = \mathcal{K}(y), \quad u' = \mathcal{K}(y + \textcolor{red}{v})$$

Lipschitz-Bounded Policy Networks

Construct a new *policy class*:

- ✓ 1. Robust to perturbations by construction
- ✓ 2. Differentiable, scalable parametrisation
- ✓ 3. Expressive parametrisation

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools



Lipschitz-Bounded Policy Network:
 $\|\Delta u\|_2 \leq \gamma \|\Delta y\|_2$

$$u = \mathcal{K}(y), \quad u' = \mathcal{K}(y + \textcolor{red}{v})$$

$$\|u' - u\|_2 \leq \gamma \|\textcolor{red}{v}\|_2$$

Playing Atari Pong (Noise-Free)

[15] Barbara, Wang, & Manchester, On Robust Reinforcement Learning with Lipschitz-Bounded Policy Networks, *SysDO* (2024).

The University of Sydney

Playing Atari Pong (Noise-Free)



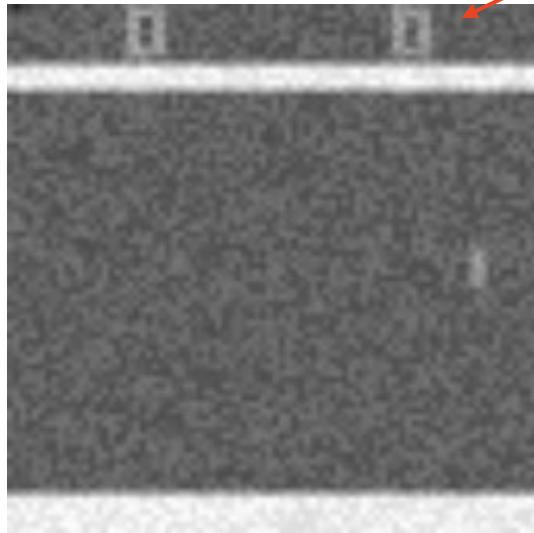
Unconstrained CNN

Trained RL Agent

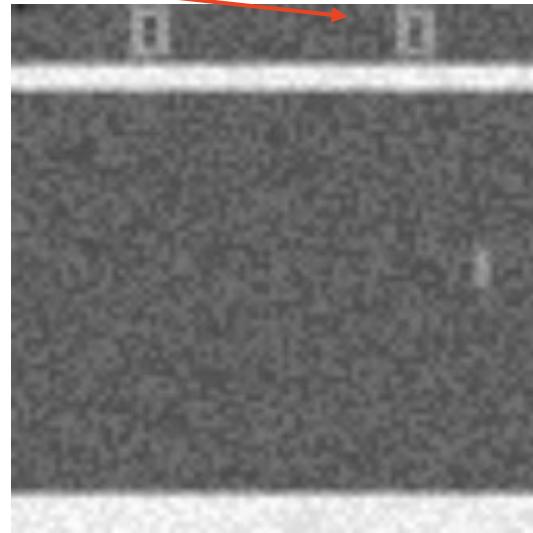


Lipschitz-Bounded CNN

Playing Atari Pong (Noise-Affected)



Unconstrained CNN



Lipschitz-Bounded CNN

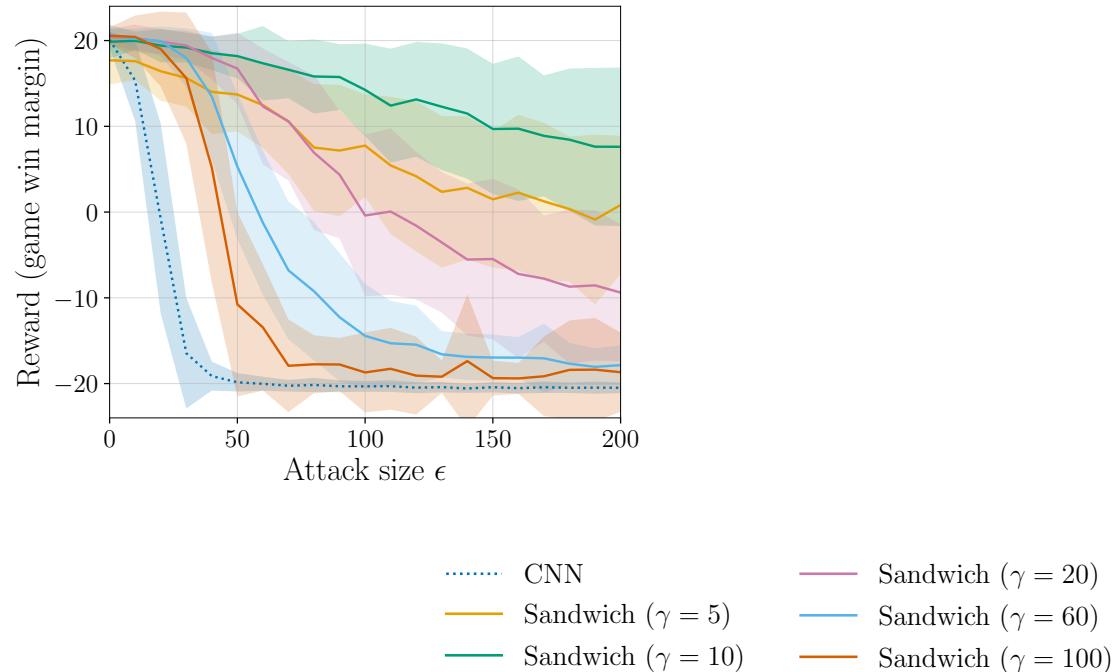
Tuneable Robustness to Adversarial Attacks

[15] Barbara, Wang, & Manchester, On Robust Reinforcement Learning with Lipschitz-Bounded Policy Networks, *SysDO* (2024).

The University of Sydney

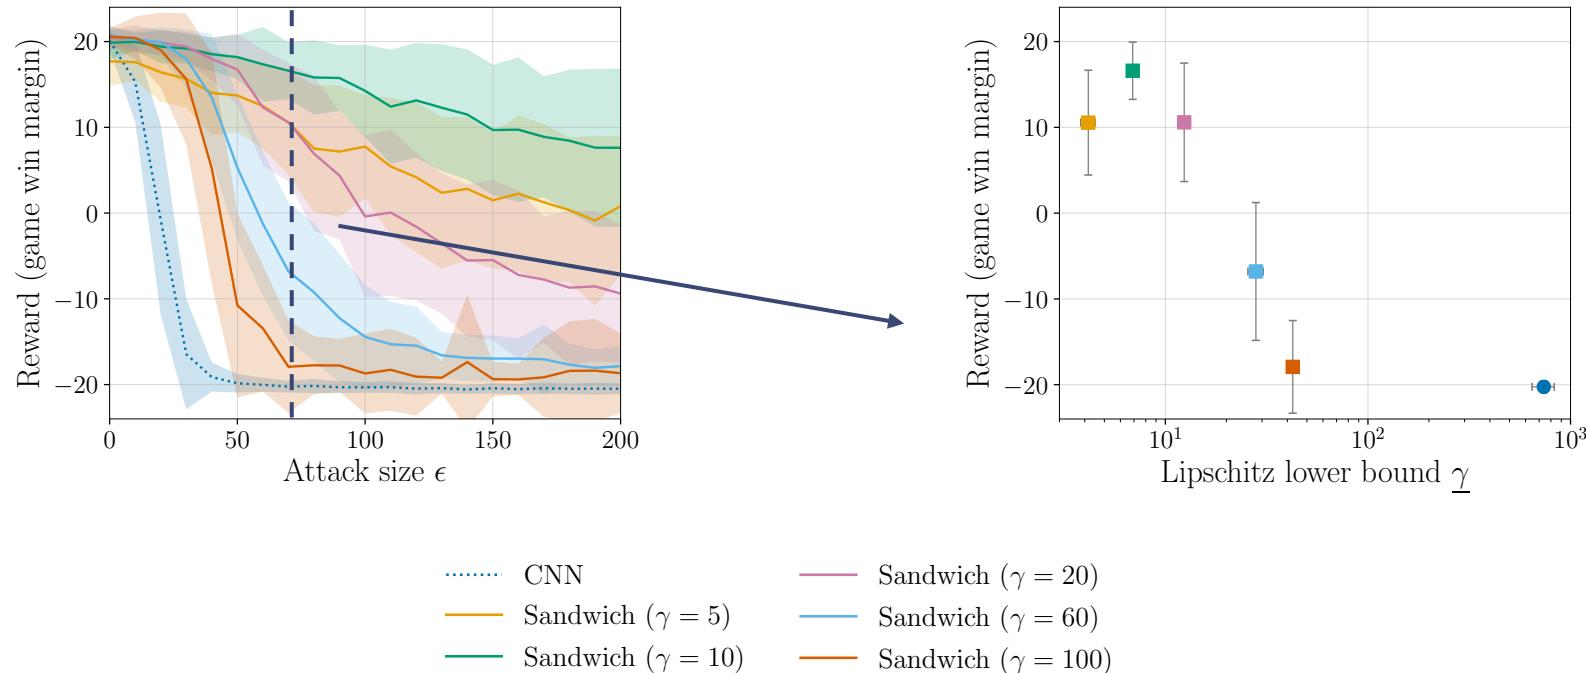
19

Tuneable Robustness to Adversarial Attacks



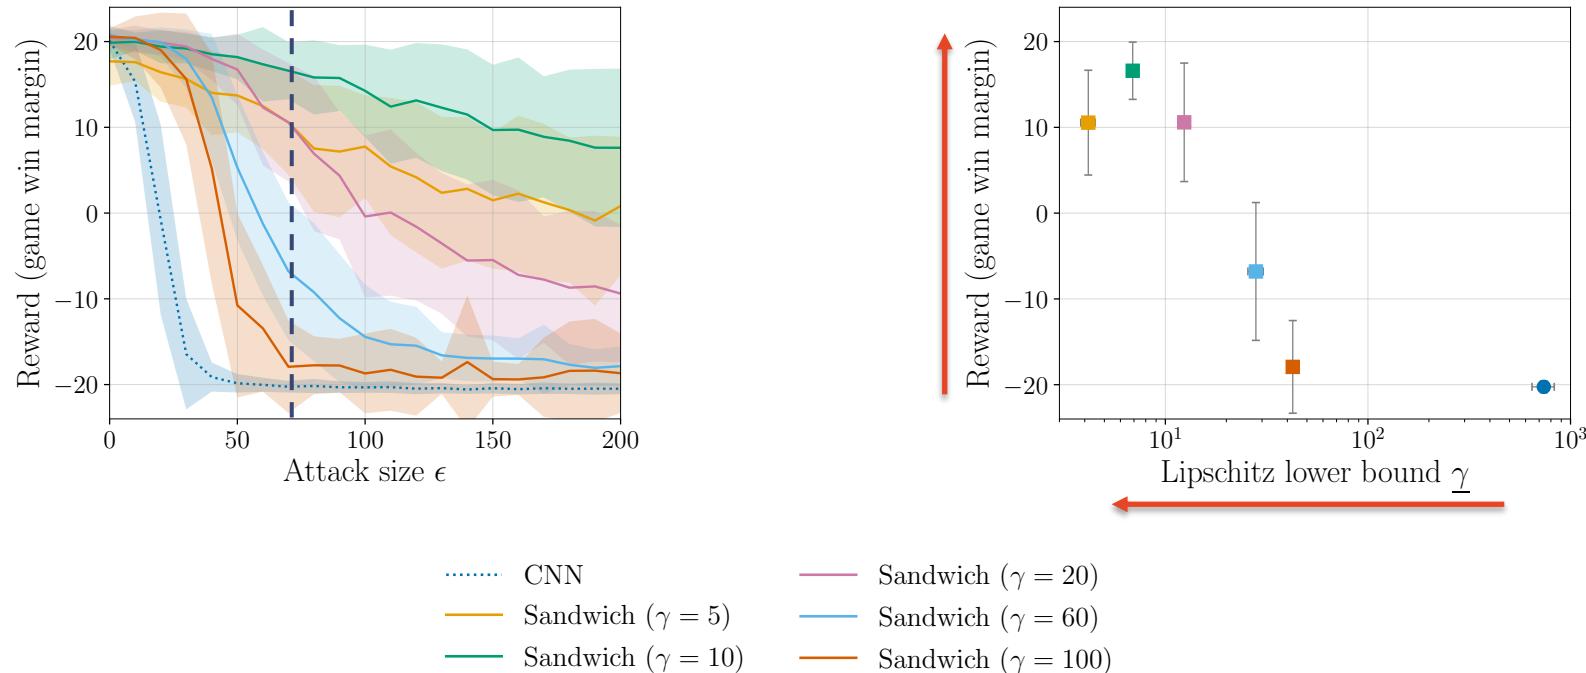
[15] Barbara, Wang, & Manchester, On Robust Reinforcement Learning with Lipschitz-Bounded Policy Networks, *SysDO* (2024).

Tuneable Robustness to Adversarial Attacks



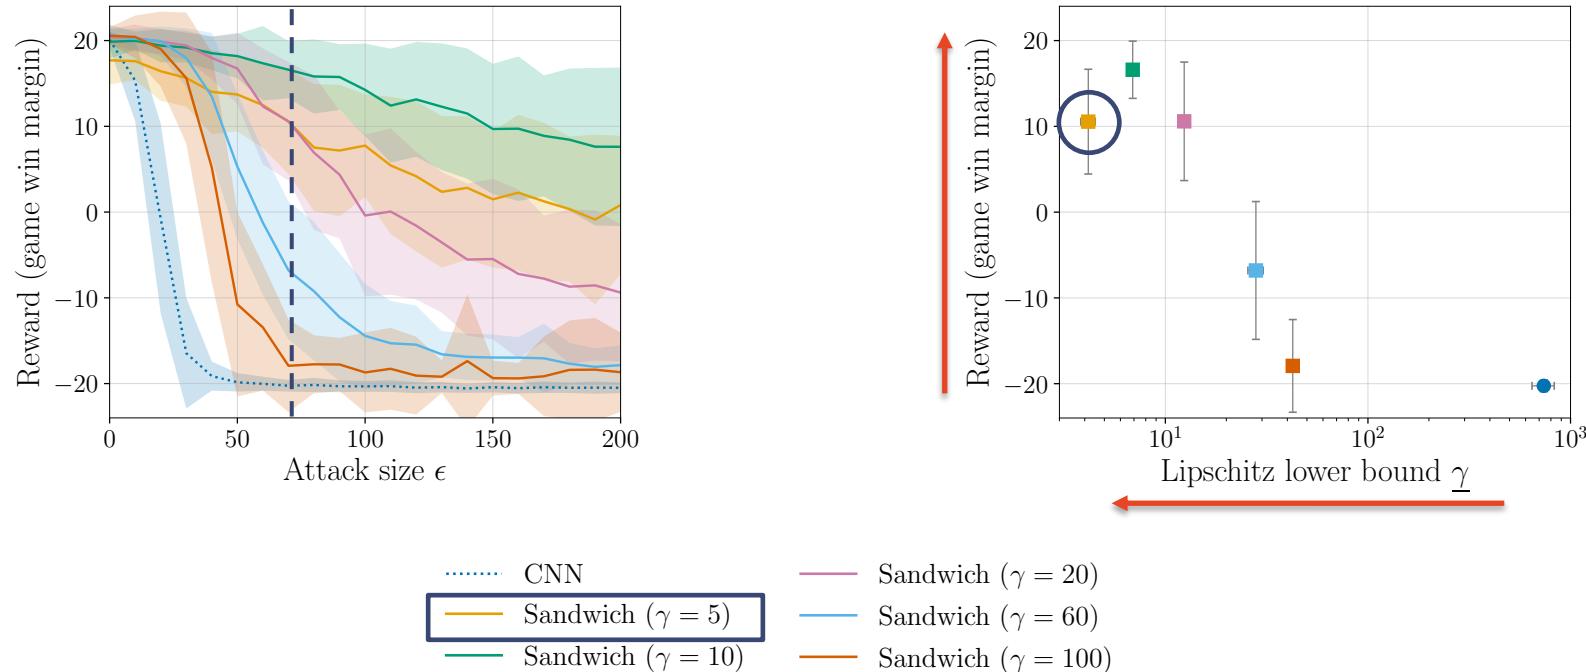
[15] Barbara, Wang, & Manchester, On Robust Reinforcement Learning with Lipschitz-Bounded Policy Networks, *SysDO* (2024).

Tuneable Robustness to Adversarial Attacks



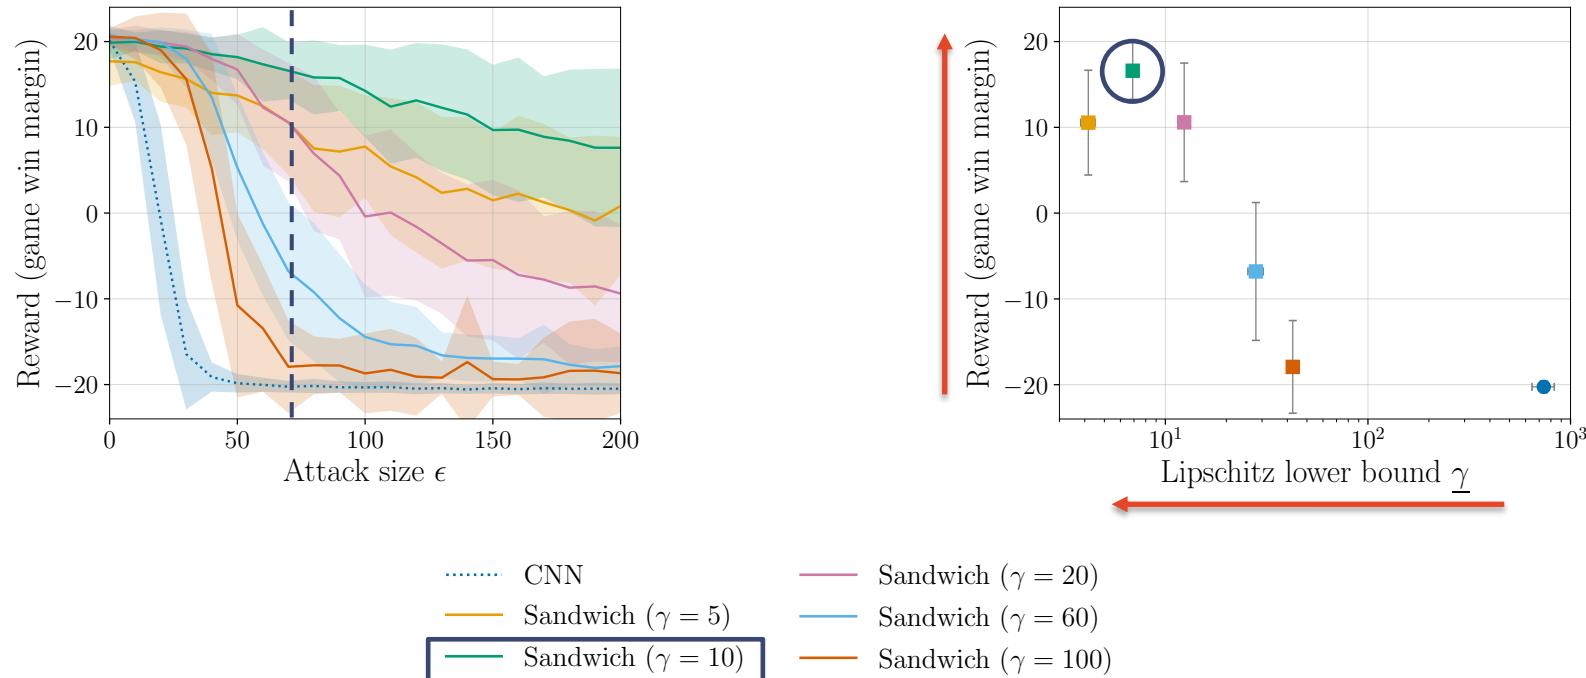
[15] Barbara, Wang, & Manchester, On Robust Reinforcement Learning with Lipschitz-Bounded Policy Networks, *SysDO* (2024).

Tuneable Robustness to Adversarial Attacks



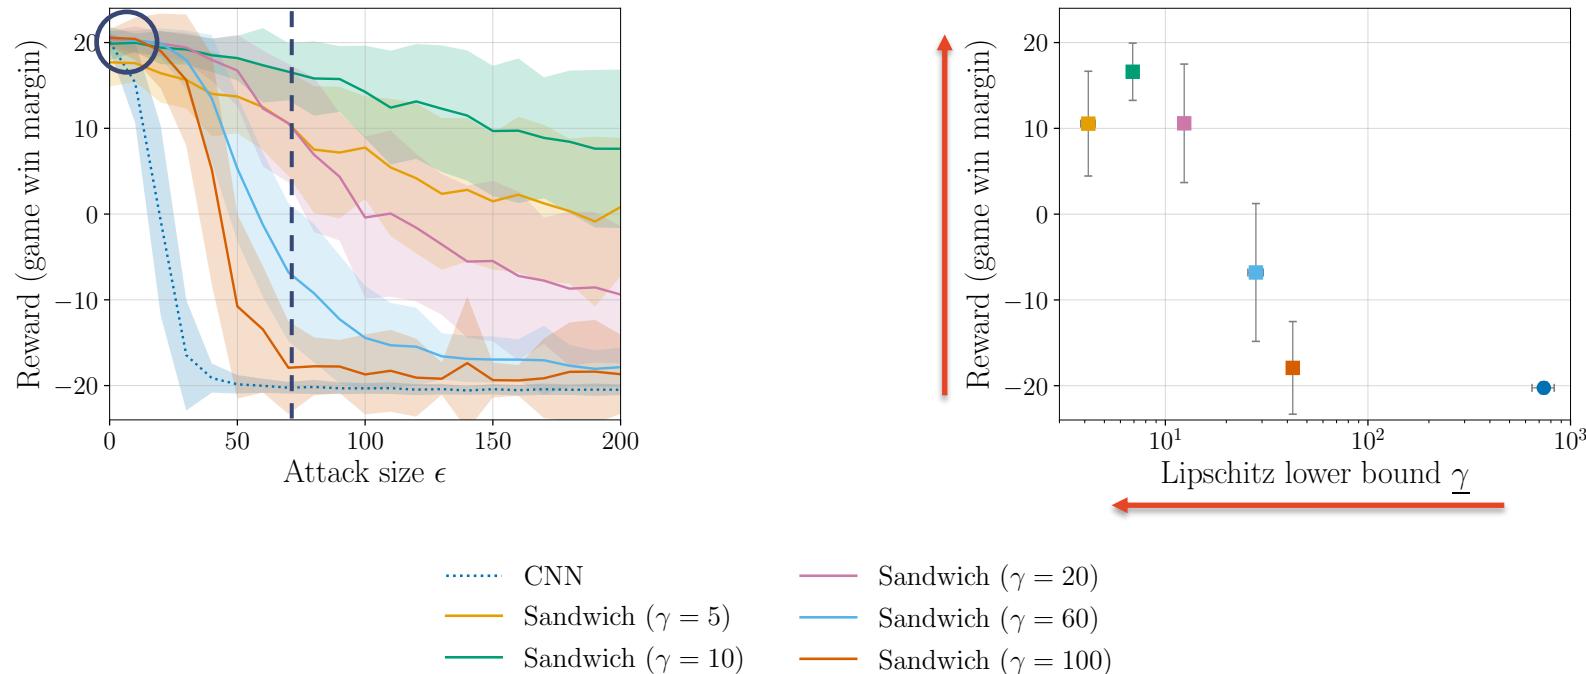
[15] Barbara, Wang, & Manchester, On Robust Reinforcement Learning with Lipschitz-Bounded Policy Networks, SysDO (2024).

Tuneable Robustness to Adversarial Attacks



[15] Barbara, Wang, & Manchester, On Robust Reinforcement Learning with Lipschitz-Bounded Policy Networks, SysDO (2024).

Tuneable Robustness to Adversarial Attacks



[15] Barbara, Wang, & Manchester, On Robust Reinforcement Learning with Lipschitz-Bounded Policy Networks, *SysDO* (2024).

Outline

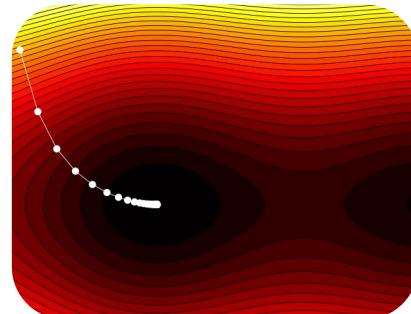
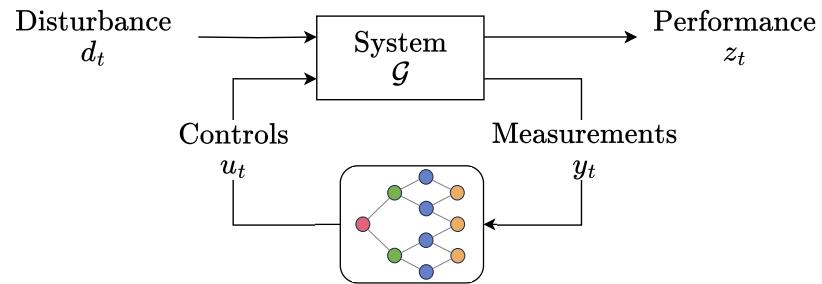
1. Motivation
2. Background: robust neural networks
3. Lipschitz-bounded policy networks
4. React to surprises with Youla-REN (Theory)
5. React to surprises with Youla-REN (Application)
6. Robust recurrent deep networks
7. Conclusions

Stable-by-Design Neural Feedback Policies

Construct a new *policy class*:

1. Closed-loop robust stability by construction
2. Differentiable, scalable parametrisation
3. Expressive parametrisation

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools

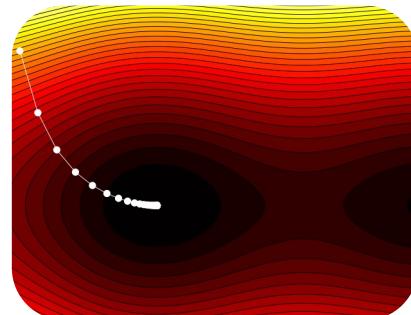
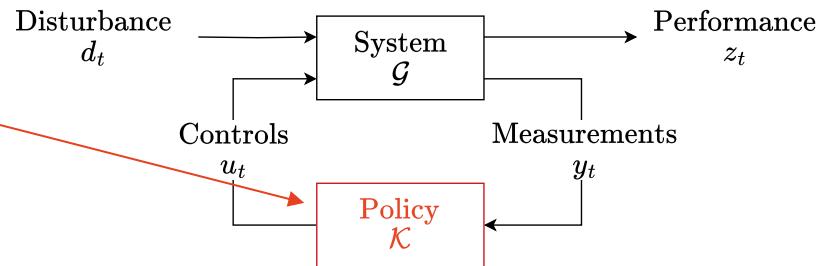


Stable-by-Design Neural Feedback Policies

Construct a new *policy class*:

1. Closed-loop robust stability by construction
2. Differentiable, scalable parametrisation
3. Expressive parametrisation

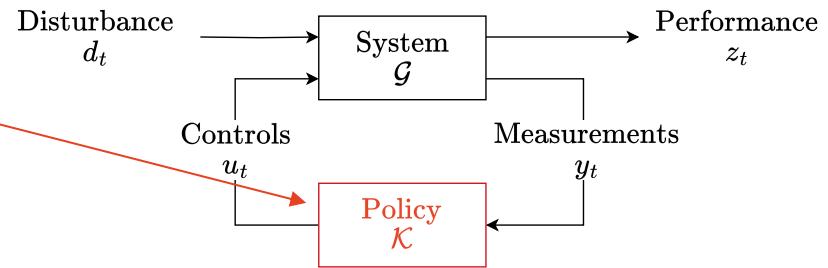
Directly learn controllers with robust stability guarantees using standard ML/deep RL tools



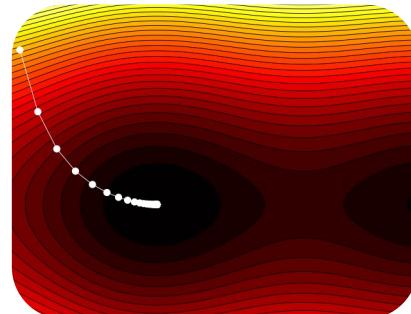
Stable-by-Design Neural Feedback Policies

Construct a new *policy class*:

- 1. Closed-loop robust stability by construction
- 2. Differentiable, scalable parametrisation
- 3. Expressive parametrisation

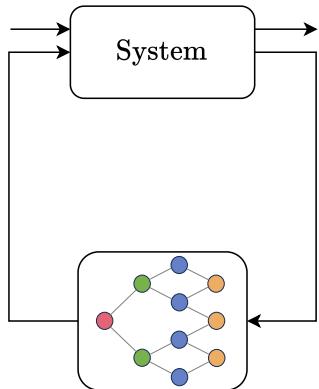


Directly learn controllers with robust stability guarantees using standard ML/deep RL tools



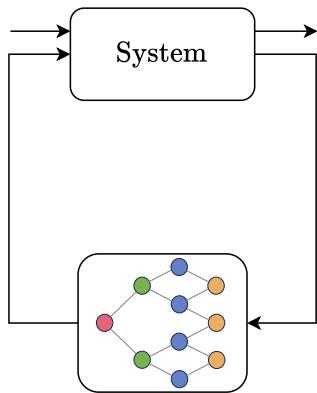
Policy Classes for Deep RL

Policy Classes for Deep RL

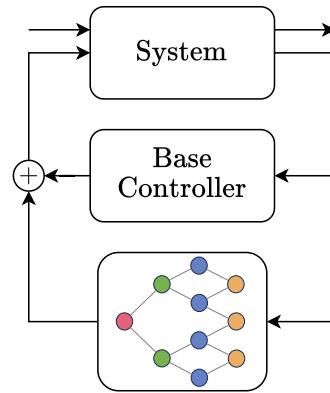


Vanilla RL

Policy Classes for Deep RL



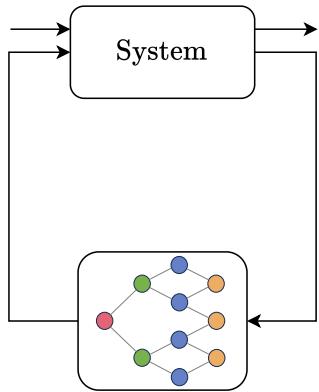
Vanilla RL



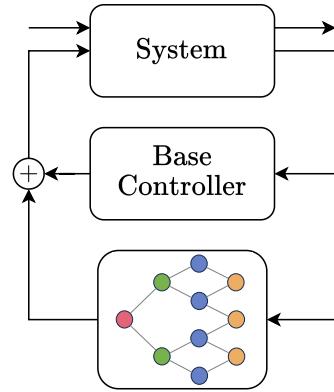
Residual RL^[16]

[16] Johannik et al., Residual Reinforcement Learning for Robot Control, ICRA (2019).

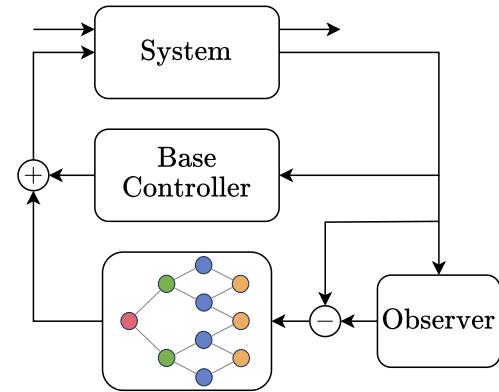
Policy Classes for Deep RL



Vanilla RL



Residual RL^[16]

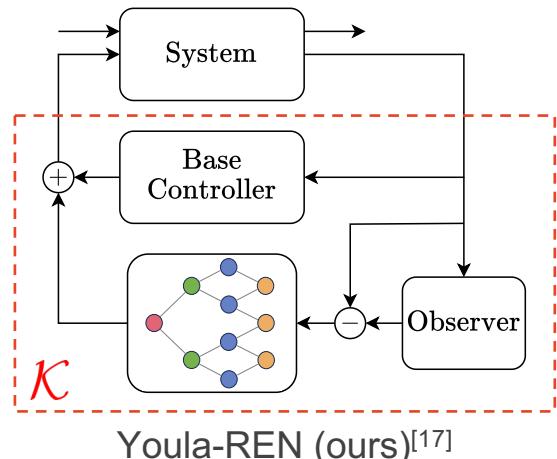
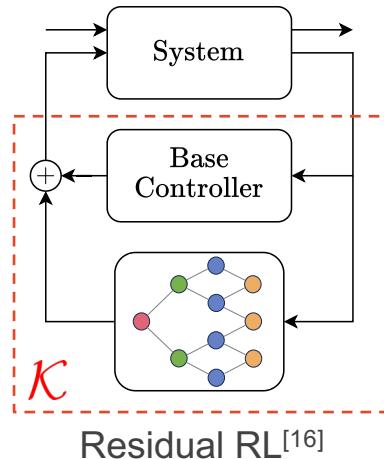
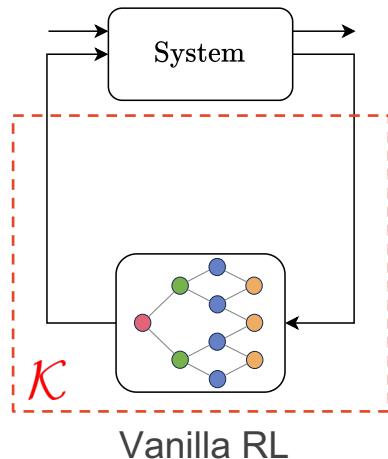


Youla-REN (ours)^[17]

[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

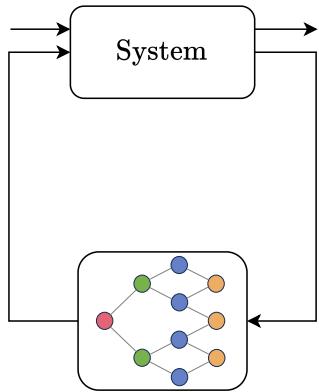
Policy Classes for Deep RL



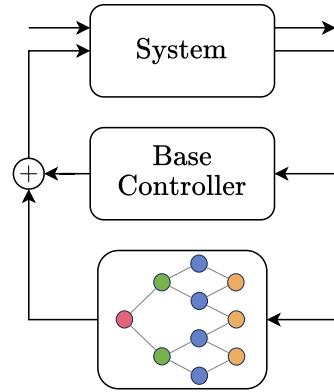
[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

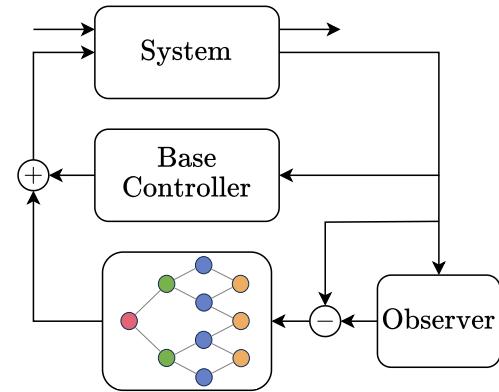
Policy Classes for Deep RL



Vanilla RL



Residual RL^[16]

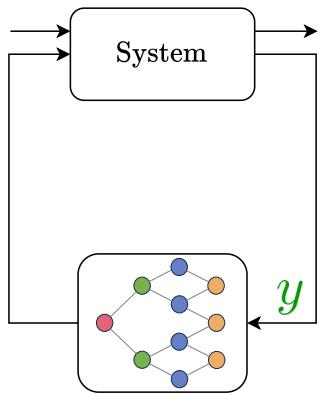


Youla-REN (ours)^[17]

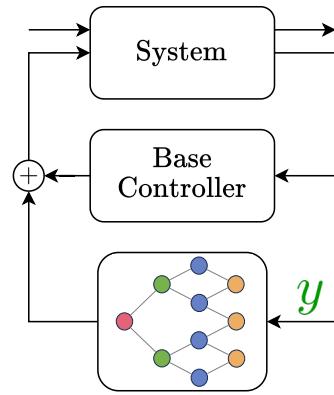
[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

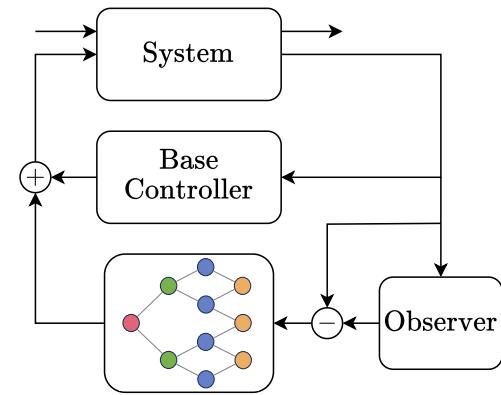
Policy Classes for Deep RL



Vanilla RL



Residual RL^[16]



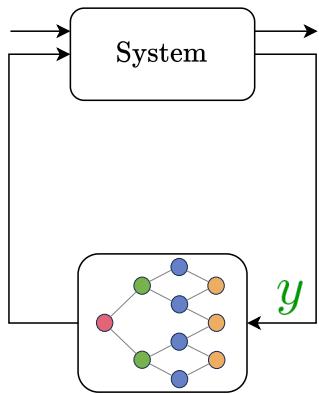
Youla-REN (ours)^[17]

Network reacts to measurements y

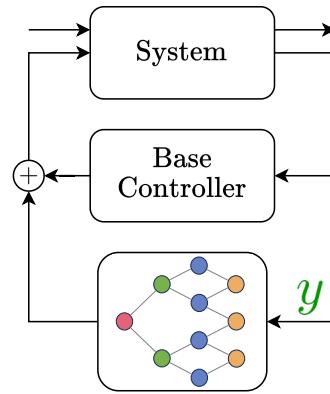
[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

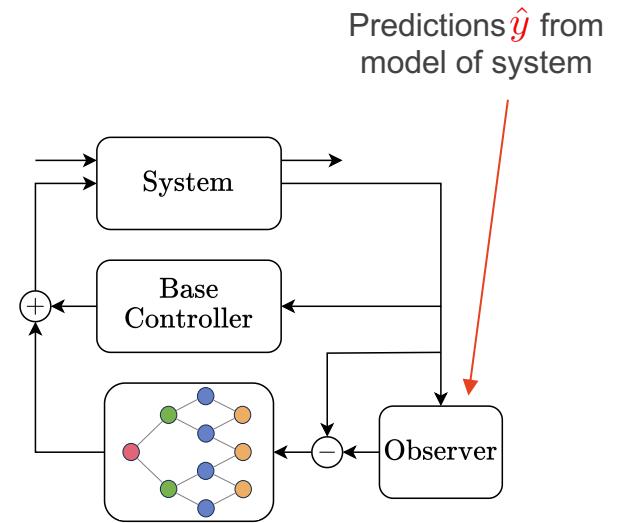
Policy Classes for Deep RL



Vanilla RL



Residual RL^[16]



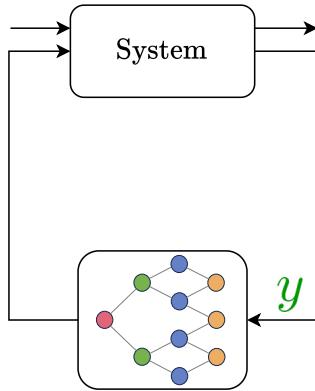
Youla-REN (ours)^[17]

Network reacts to measurements y

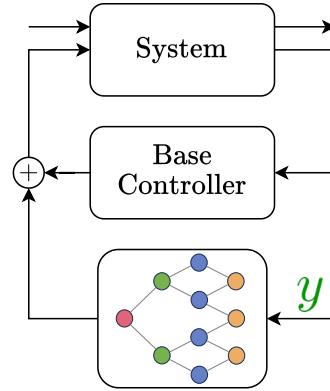
[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

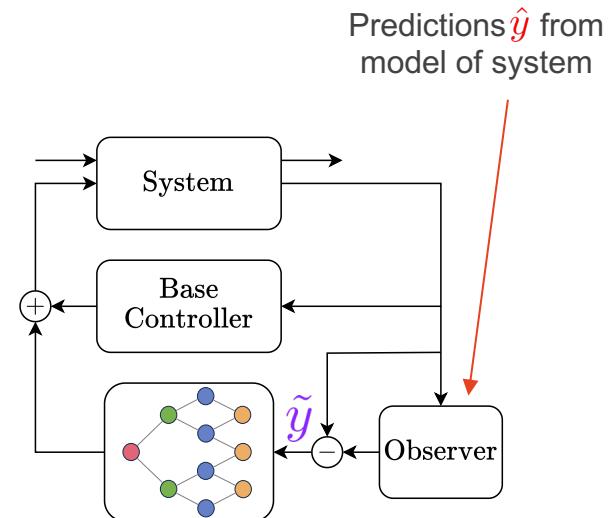
Policy Classes for Deep RL



Vanilla RL



Residual RL^[16]



Youla-REN (ours)^[17]

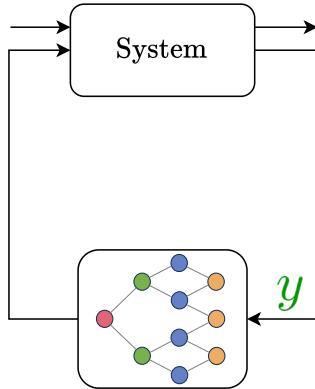
Network reacts to **measurements y**

Network reacts to “surprises” $\tilde{y} = y - \hat{y}$

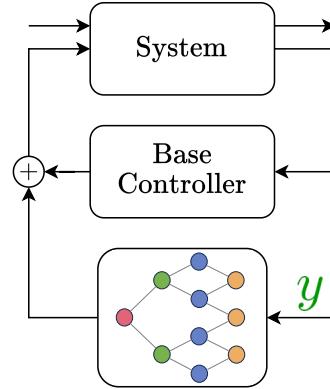
[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

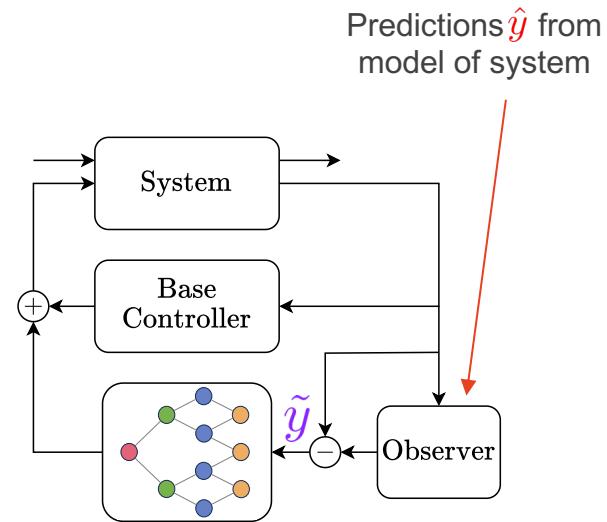
Policy Classes for Deep RL



Vanilla RL



Residual RL^[16]



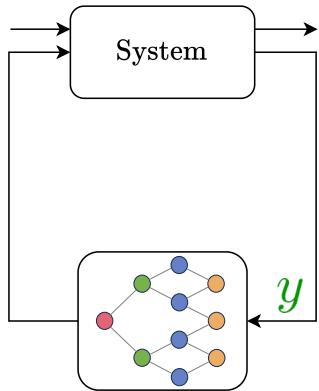
Youla-REN (ours)^[17]

- React to measurements y : closed-loop stability via small-gain theorem

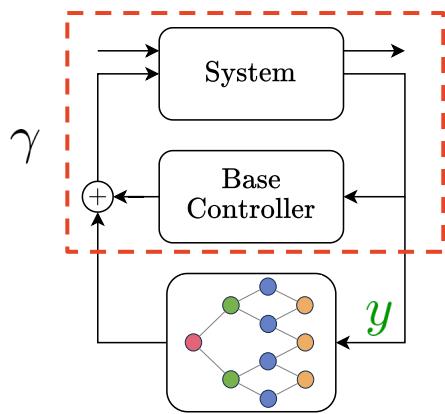
[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

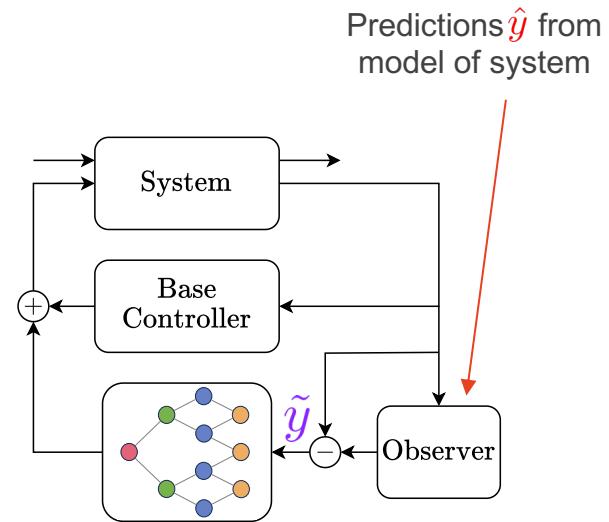
Policy Classes for Deep RL



Vanilla RL



Residual RL^[16]



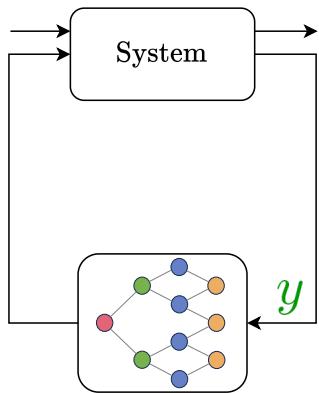
Youla-REN (ours)^[17]

- React to measurements y : closed-loop stability via small-gain theorem

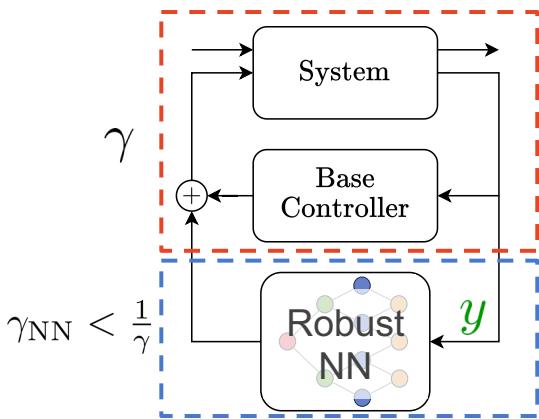
[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

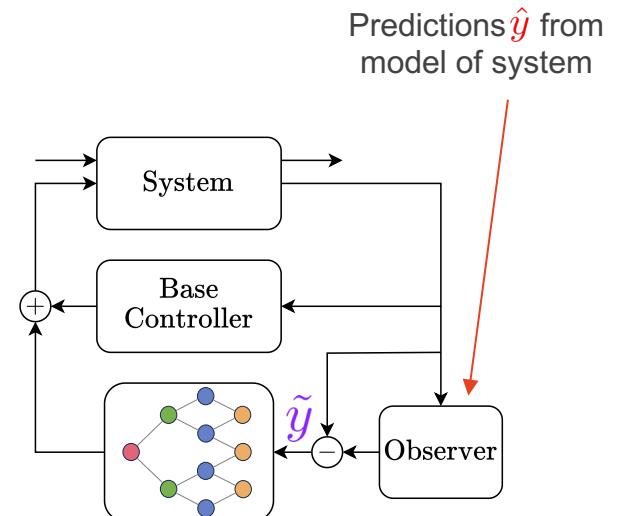
Policy Classes for Deep RL



Vanilla RL



Residual RL^[16]



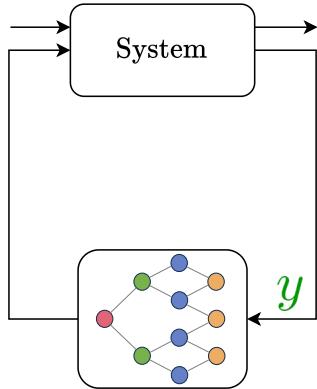
Youla-REN (ours)^[17]

- React to measurements y : closed-loop stability via small-gain theorem

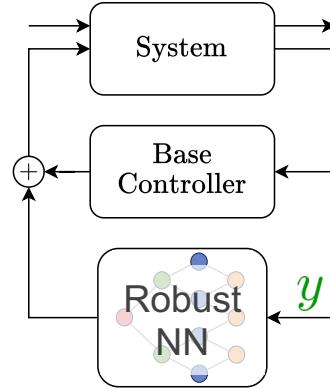
[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

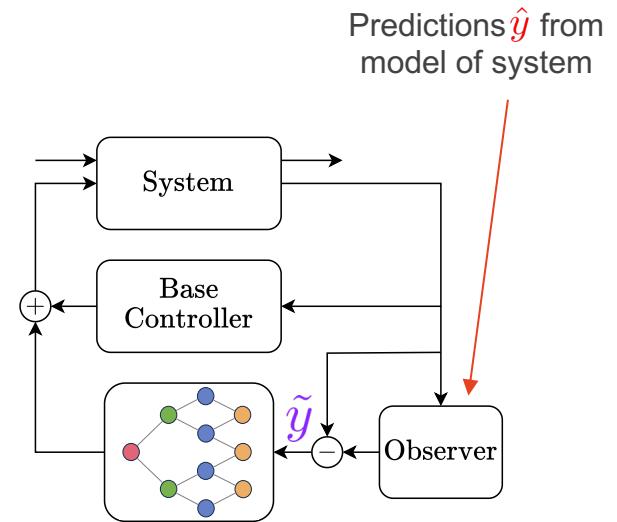
Policy Classes for Deep RL



Vanilla RL



Residual RL^[16]



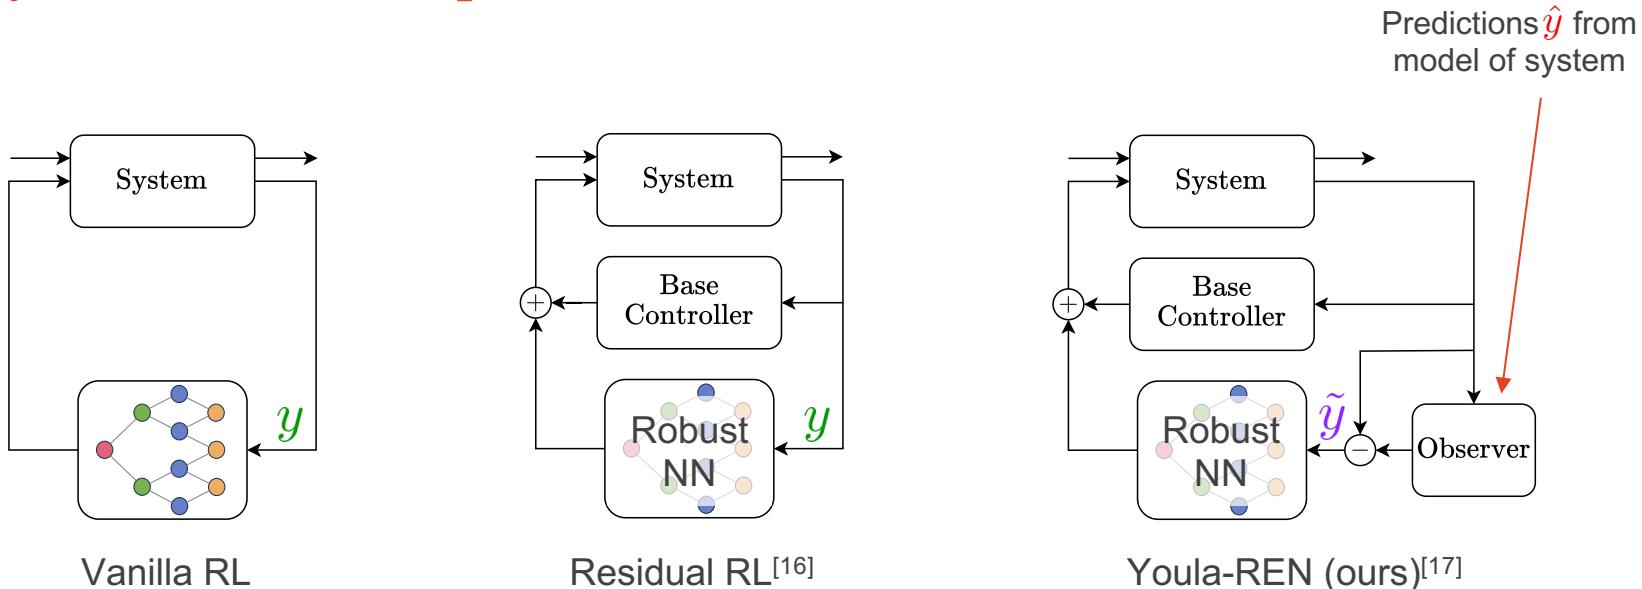
Youla-REN (ours)^[17]

- React to measurements y : closed-loop stability via small-gain theorem

[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

Policy Classes for Deep RL

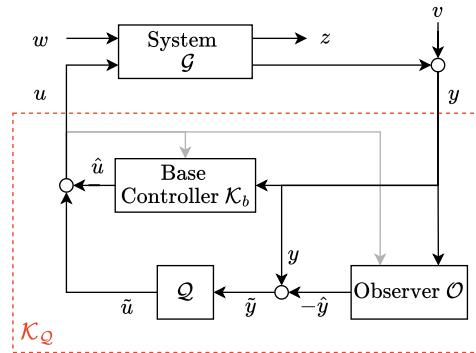


- React to **measurements y** : closed-loop stability via small-gain theorem
- React to **surprises \tilde{y}** : closed-loop stability **without** gain restrictions (under assumptions)

[16] Johannik et al., Residual Reinforcement Learning for Robot Control, *ICRA* (2019).

[17] Barbara, Wang, & Manchester, React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, *arXiv* (2025).

Youla Parametrisation Guarantees Closed-Loop Stability

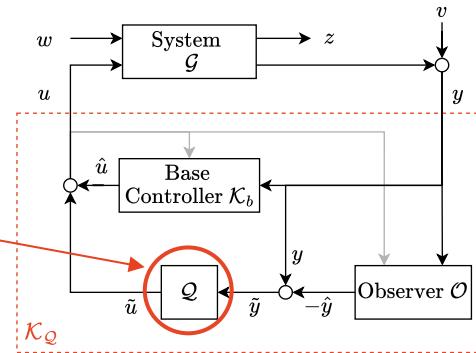


Youla Parametrisation Guarantees Closed-Loop Stability

“Youla Parameter” Q

Nonlinear (possibly dynamic) system.

Eg: neural network, or more general system



Youla Parametrisation Guarantees Closed-Loop Stability

Theorem. Consider the LTI system

$$\mathcal{G} : \begin{cases} x_{t+1} = Ax_t + Bu_t + w_t \\ y_t = Cx_t + v_t \end{cases}$$

together with an observer and base controller

$$\mathcal{O}, \mathcal{K}_b : \begin{cases} \hat{x}_{t+1} = A\hat{x}_t + Bu_t + L(y - C\hat{x}_t) \\ \hat{u}_t = -K\hat{x}_t \end{cases}$$

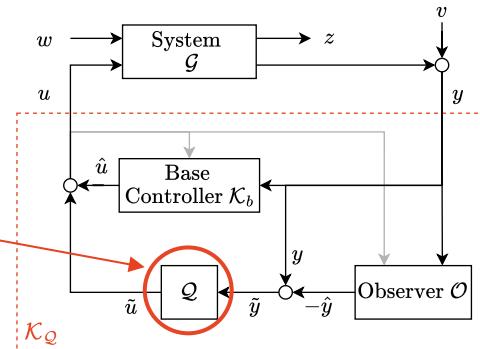
designed such that $(A - BK)$ and $(A - LC)$ are stable. Writing $d = [w; v]$ and $z = [x; u]$, then the following statement is true for the closed-loop system under the Youla controller $u = \hat{u} + \mathcal{Q}(y - C\hat{x})$:

A controller $u = \mathcal{K}(y)$ with a valid state-space representation can achieve a contracting and Lipschitz closed loop $d \mapsto z$ if and only if it can be written in the form of the Youla controller with a contracting and Lipschitz \mathcal{Q} .

“Youla Parameter” \mathcal{Q}

Nonlinear (possibly dynamic) system.

Eg: neural network, or more general system



Youla Parametrisation Guarantees Closed-Loop Stability

Theorem. Consider the LTI system

$$\mathcal{G} : \begin{cases} x_{t+1} = Ax_t + Bu_t + w_t \\ y_t = Cx_t + v_t \end{cases}$$

together with an observer and base controller

$$\mathcal{O}, \mathcal{K}_b : \begin{cases} \hat{x}_{t+1} = A\hat{x}_t + Bu_t + L(y - C\hat{x}_t) \\ \hat{u}_t = -K\hat{x}_t \end{cases}$$

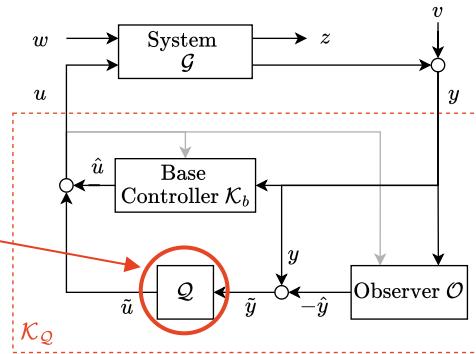
designed such that $(A - BK)$ and $(A - LC)$ are stable. Writing $d = [w; v]$ and $z = [x; u]$, then the following statement is true for the closed-loop system under the Youla controller $u = \hat{u} + Q(y - C\hat{x})$:

A controller $u = \mathcal{K}(y)$ with a valid state-space representation can achieve a contracting and Lipschitz closed loop $d \mapsto z$ if and only if it can be written in the form of the Youla controller with a contracting and Lipschitz Q .

“Youla Parameter” Q

Nonlinear (possibly dynamic) system.

Eg: neural network, or more general system



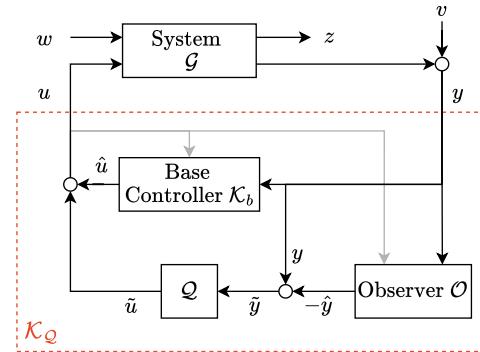
What it means:

- Stabilising controller \iff Youla controller with a **stable (contracting & Lipschitz) Q**
 - No gain restrictions on Q
 - For linear^[18] and nonlinear^[19] controllers
- I.e., the **policy class \mathcal{K}_Q** covers **all and only** stabilising controllers for LTI systems!

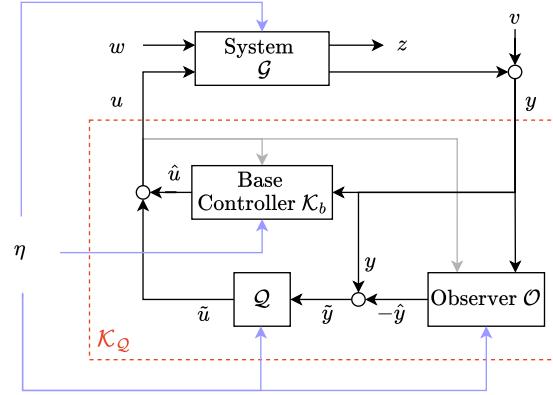
[18] Doyle, Matrix Interpolation Theory and Optimal Control, *PhD Thesis, Berkley* (1984).

[19] Wang, Barbara, Revay, & Manchester, Learning Over All Stabilizing Nonlinear Controllers for a Partially-Observed Linear System, *L-CSS* (2022).

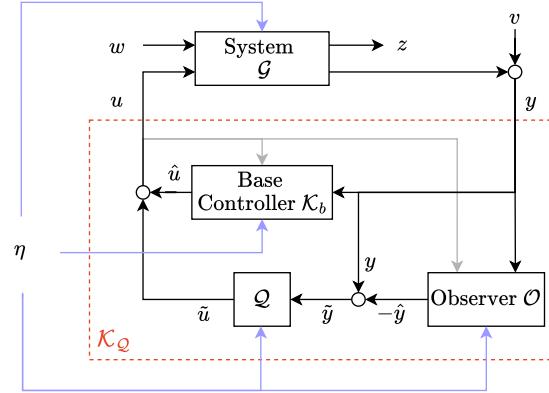
Contracting & Lipschitz Closed Loops for Nonlinear Systems



Contracting & Lipschitz Closed Loops for Nonlinear Systems



Contracting & Lipschitz Closed Loops for Nonlinear Systems



$$\mathcal{G} : \begin{cases} x_{t+1} = f(x_t, \eta_t, u_t) + w_t \\ y_t = h(x_t) + v_t \end{cases}$$

$$\mathcal{K}_b : \begin{cases} s_{t+1} = f_o(s_t, \eta_t, u_t, y_t) \\ \hat{u}_t = k(s_t, \eta_t, y_t) \end{cases}$$

$$\mathcal{O} : \begin{cases} \hat{x}_{t+1} = f_o(\hat{x}_t, \eta_t, u_t, y_t) \\ \hat{y}_t = h(\hat{x}_t) \end{cases}$$

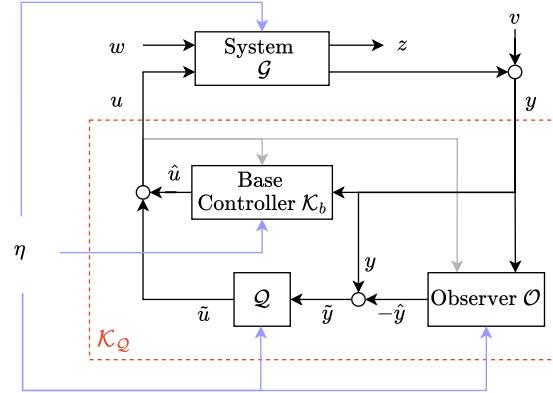
$$\mathcal{Q} : \begin{cases} q_{t+1} = f_q(q_t, \eta_t, \tilde{y}_t) \\ \tilde{u}_t = h_q(q_t, \eta_t, \tilde{y}_t) \end{cases}$$

Contracting & Lipschitz Closed Loops for Nonlinear Systems

Theorem. Consider the partially-observed, nonlinear system \mathcal{G} together with the base controller and observer \mathcal{K}_b and \mathcal{O} (respectively) which satisfy the following assumptions, where $d = [w; v]$ and $z = [x; u]$:

- \mathcal{K}_b achieves a contracting & Lipschitz closed loop $(\eta, d, \tilde{u}) \mapsto z$ with \mathcal{G} .
- \mathcal{O} is contracting and the map $(\eta, u, y) \mapsto \hat{x}$ is Lipschitz.
- Observer correctness: $f(x_t, \eta_t, u_t) = f_o(s_t, \eta_t, u_t, h(x_t))$.
- All systems have piecewise-diff'ble dynamics and Lipschitz output maps.

Then, the following statements are true for the closed-loop system under the Youla controller $u = \hat{u} + \mathcal{Q}(y - \hat{y})$ with a **contracting and Lipschitz** \mathcal{Q} :



$$\mathcal{G} : \begin{cases} x_{t+1} = f(x_t, \eta_t, u_t) + w_t \\ y_t = h(x_t) + v_t \end{cases}$$

$$\mathcal{K}_b : \begin{cases} s_{t+1} = f_o(s_t, \eta_t, u_t, y_t) \\ \hat{u}_t = k(s_t, \eta_t, y_t) \end{cases}$$

$$\mathcal{O} : \begin{cases} \hat{x}_{t+1} = f_o(\hat{x}_t, \eta_t, u_t, y_t) \\ \hat{y}_t = h(\hat{x}_t) \end{cases}$$

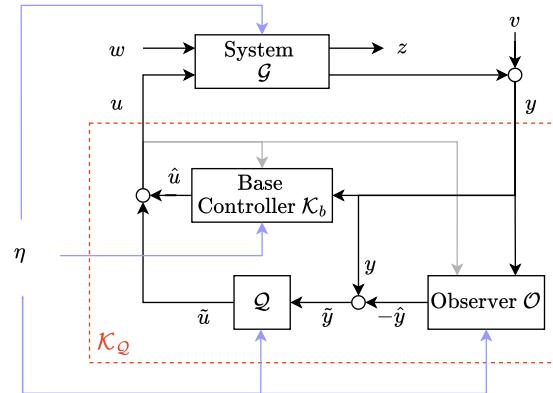
$$\mathcal{Q} : \begin{cases} q_{t+1} = f_q(q_t, \eta_t, \tilde{y}_t) \\ \tilde{u}_t = h_q(q_t, \eta_t, \tilde{y}_t) \end{cases}$$

Contracting & Lipschitz Closed Loops for Nonlinear Systems

Theorem. Consider the partially-observed, nonlinear system \mathcal{G} together with the base controller and observer \mathcal{K}_b and \mathcal{O} (respectively) which satisfy the following assumptions, where $d = [w; v]$ and $z = [x; u]$:

- \mathcal{K}_b achieves a contracting & Lipschitz closed loop $(\eta, d, \tilde{u}) \mapsto z$ with \mathcal{G} .
- \mathcal{O} is contracting and the map $(\eta, u, y) \mapsto \hat{x}$ is Lipschitz.
- Observer correctness: $f(x_t, \eta_t, u_t) = f_o(s_t, \eta_t, u_t, h(x_t))$.
- All systems have piecewise-diff'ble dynamics and Lipschitz output maps.

Then, the following statements are true for the closed-loop system under the Youla controller $u = \hat{u} + \mathcal{Q}(y - \hat{y})$ with a **contracting and Lipschitz** \mathcal{Q} :



$$\mathcal{G} : \begin{cases} x_{t+1} = f(x_t, \eta_t, u_t) + w_t \\ y_t = h(x_t) + v_t \end{cases}$$

$$\mathcal{K}_b : \begin{cases} s_{t+1} = f_o(s_t, \eta_t, u_t, y_t) \\ \hat{u}_t = k(s_t, \eta_t, y_t) \end{cases}$$

$$\mathcal{O} : \begin{cases} \hat{x}_{t+1} = f_o(\hat{x}_t, \eta_t, u_t, y_t) \\ \hat{y}_t = h(\hat{x}_t) \end{cases}$$

$$\mathcal{Q} : \begin{cases} q_{t+1} = f_q(q_t, \eta_t, \tilde{y}_t) \\ \tilde{u}_t = h_q(q_t, \eta_t, \tilde{y}_t) \end{cases}$$

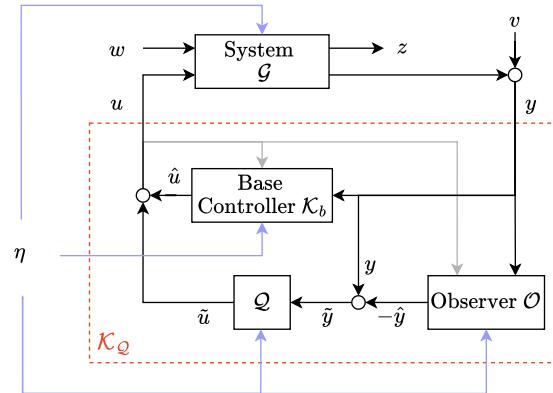
Contracting & Lipschitz Closed Loops for Nonlinear Systems

Theorem. Consider the partially-observed, nonlinear system \mathcal{G} together with the base controller and observer \mathcal{K}_b and \mathcal{O} (respectively) which satisfy the following assumptions, where $d = [w; v]$ and $z = [x; u]$:

- \mathcal{K}_b achieves a contracting & Lipschitz closed loop $(\eta, d, \tilde{u}) \mapsto z$ with \mathcal{G} .
- \mathcal{O} is contracting and the map $(\eta, u, y) \mapsto \hat{x}$ is Lipschitz.
- Observer correctness: $f(x_t, \eta_t, u_t) = f_o(s_t, \eta_t, u_t, h(x_t))$.
- All systems have piecewise-diff'ble dynamics and Lipschitz output maps.

Then, the following statements are true for the closed-loop system under the Youla controller $u = \hat{u} + \mathcal{Q}(y - \hat{y})$ **with a contracting and Lipschitz \mathcal{Q}** :

1. With zero disturbances (i.e., $d \equiv 0$, $x_0 = \hat{x}_0$), the closed-loop system is contracting and $\eta \mapsto z$ is Lipschitz.



$$\mathcal{G} : \begin{cases} x_{t+1} = f(x_t, \eta_t, u_t) + w_t \\ y_t = h(x_t) + v_t \end{cases}$$

$$\mathcal{K}_b : \begin{cases} s_{t+1} = f_o(s_t, \eta_t, u_t, y_t) \\ \hat{u}_t = k(s_t, \eta_t, y_t) \end{cases}$$

$$\mathcal{O} : \begin{cases} \hat{x}_{t+1} = f_o(\hat{x}_t, \eta_t, u_t, y_t) \\ \hat{y}_t = h(\hat{x}_t) \end{cases}$$

$$\mathcal{Q} : \begin{cases} q_{t+1} = f_q(q_t, \eta_t, \tilde{y}_t) \\ \tilde{u}_t = h_q(q_t, \eta_t, \tilde{y}_t) \end{cases}$$

Contracting & Lipschitz Closed Loops for Nonlinear Systems

Theorem. Consider the partially-observed, nonlinear system \mathcal{G} together with the base controller and observer \mathcal{K}_b and \mathcal{O} (respectively) which satisfy the following assumptions, where $d = [w; v]$ and $z = [x; u]$:

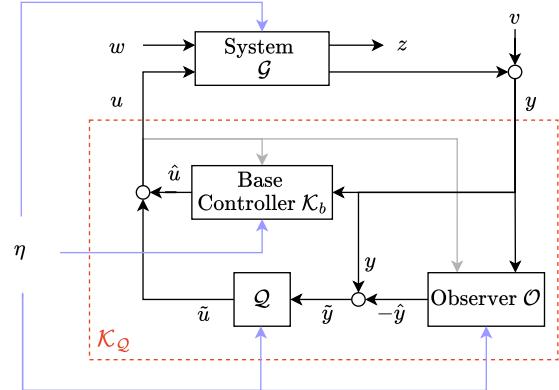
- \mathcal{K}_b achieves a contracting & Lipschitz closed loop $(\eta, d, \tilde{u}) \mapsto z$ with \mathcal{G} .
- \mathcal{O} is contracting and the map $(\eta, u, y) \mapsto \hat{x}$ is Lipschitz.
- Observer correctness: $f(x_t, \eta_t, u_t) = f_o(s_t, \eta_t, u_t, h(x_t))$.
- All systems have piecewise-diff'ble dynamics and Lipschitz output maps.

Then, the following statements are true for the closed-loop system under the Youla controller $u = \hat{u} + \mathcal{Q}(y - \hat{y})$ **with a contracting and Lipschitz \mathcal{Q}** :

1. With zero disturbances (i.e., $d \equiv 0$, $x_0 = \hat{x}_0$), the closed-loop system is contracting and $\eta \mapsto z$ is Lipschitz.
2. Let \star denote any zero-disturbance closed-loop trajectory. Then with non-zero disturbances (i.e., $d \not\equiv 0$, $x_0 \neq \hat{x}_0$), the closed-loop system satisfies

$$\|z - z^\star\|_T \leq \gamma(\|\eta - \eta^\star\|_T + \|d\|_T) + \kappa(\bar{x}_0, \bar{x}_0^\star)$$

for all $T \in \mathbb{N}$, \bar{x}_0, η , where $\gamma \in \mathbb{R}^+$, $\kappa(\bar{x}_0, \bar{x}_0^\star) \geq 0$, $\kappa(\bar{x}_0^\star, \bar{x}_0^\star) = 0$. We call this property **d-tube contracting and Lipschitz**.



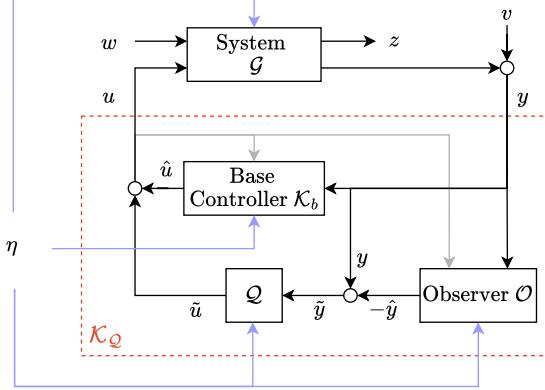
$$\mathcal{G} : \begin{cases} x_{t+1} = f(x_t, \eta_t, u_t) + w_t \\ y_t = h(x_t) + v_t \end{cases}$$

$$\mathcal{K}_b : \begin{cases} s_{t+1} = f_o(s_t, \eta_t, u_t, y_t) \\ \hat{u}_t = k(s_t, \eta_t, y_t) \end{cases}$$

$$\mathcal{O} : \begin{cases} \hat{x}_{t+1} = f_o(\hat{x}_t, \eta_t, u_t, y_t) \\ \hat{y}_t = h(\hat{x}_t) \end{cases}$$

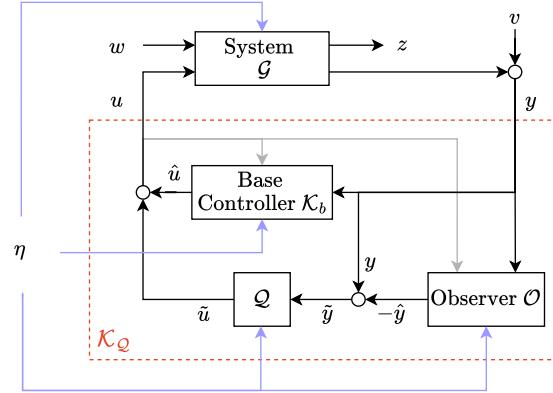
$$\mathcal{Q} : \begin{cases} q_{t+1} = f_q(q_t, \eta_t, \tilde{y}_t) \\ \tilde{u}_t = h_q(q_t, \eta_t, \tilde{y}_t) \end{cases}$$

d-Tube: What it means



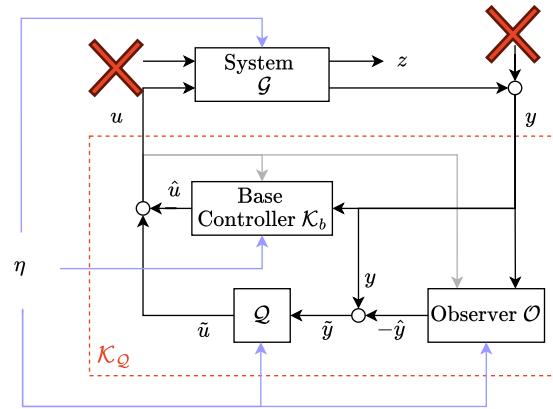
d-Tube: What it means

- Known inputs: the closed-loop system of \mathcal{K}_Q and \mathcal{G} is contracting & Lipschitz



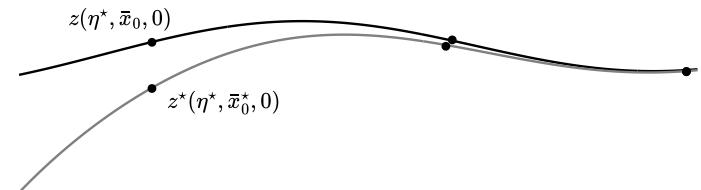
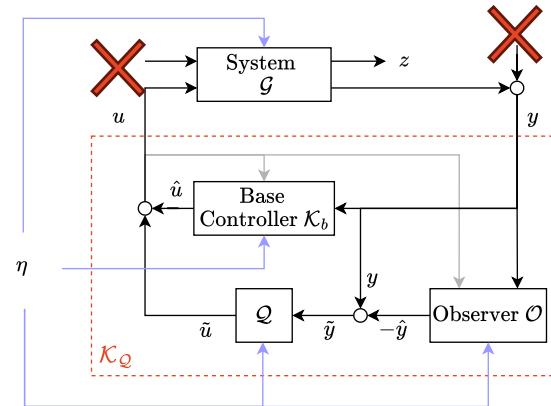
d-Tube: What it means

- Known inputs: the closed-loop system of \mathcal{K}_Q and \mathcal{G} is contracting & Lipschitz



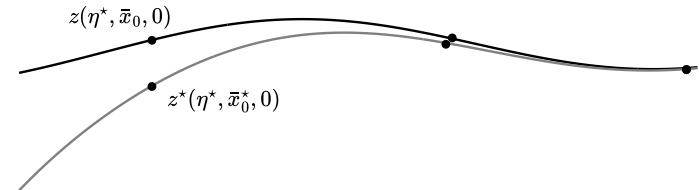
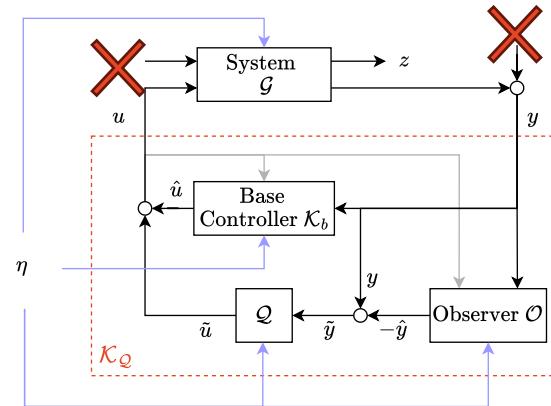
d-Tube: What it means

- Known inputs: the closed-loop system of \mathcal{K}_Q and \mathcal{G} is contracting & Lipschitz



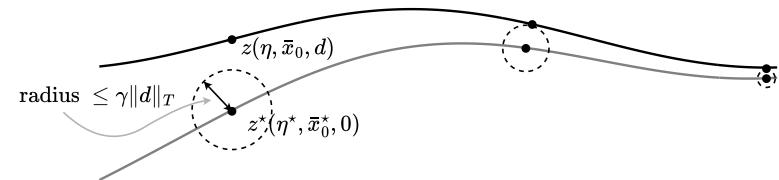
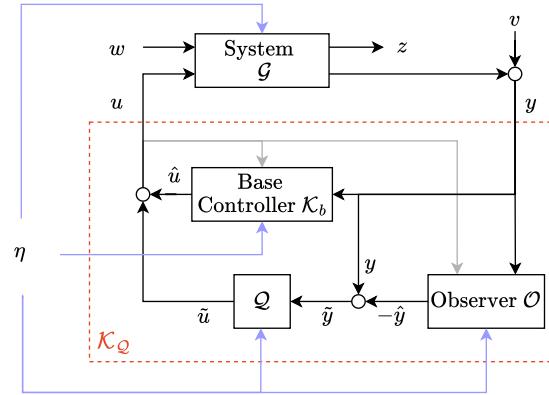
d-Tube: What it means

- Known inputs: the closed-loop system of \mathcal{K}_Q and \mathcal{G} is contracting & Lipschitz
- Unknown inputs: trajectories converge to a tube around the disturbance-free trajectories.



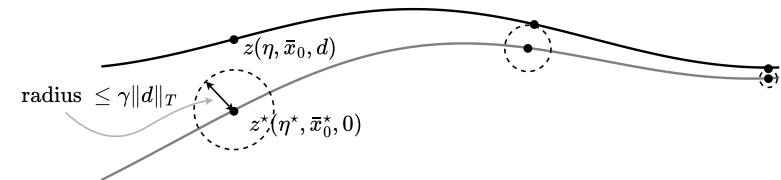
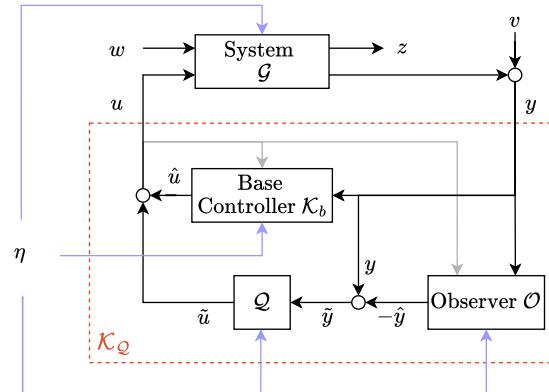
d-Tube: What it means

- Known inputs: the closed-loop system of \mathcal{K}_Q and \mathcal{G} is contracting & Lipschitz
- Unknown inputs: trajectories converge to a tube around the disturbance-free trajectories.



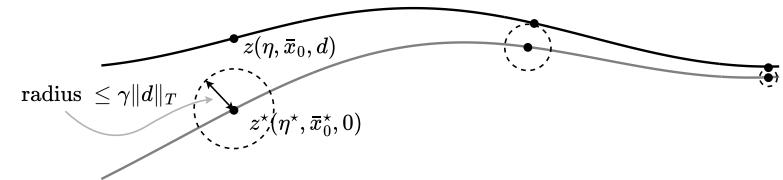
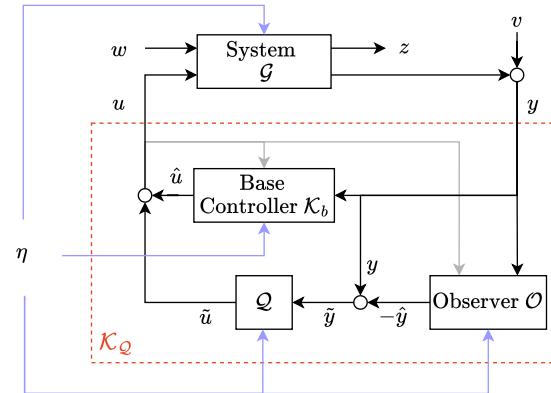
d-Tube: What it means

- Known inputs: the closed-loop system of \mathcal{K}_Q and \mathcal{G} is contracting & Lipschitz
- Unknown inputs: trajectories converge to a tube around the disturbance-free trajectories.
- Is this useful?
 - Yes – practical stability
 - Smooth response to **known** inputs
 - Bounded response to bounded **unknown** inputs



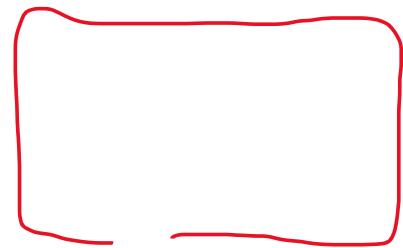
d-Tube: What it means

- Known inputs: the closed-loop system of \mathcal{K}_Q and \mathcal{G} is contracting & Lipschitz
- Unknown inputs: trajectories converge to a tube around the disturbance-free trajectories.
- Is this useful?
 - Yes – practical stability
 - Smooth response to **known** inputs
 - Bounded response to bounded **unknown** inputs
- Converse results: see thesis or paper^[17] ☺



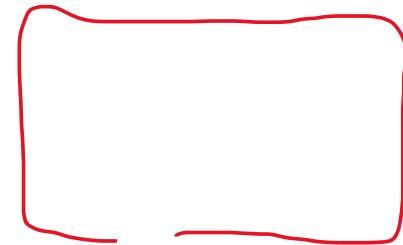
[17] Barbara et al., React to Surprises: Stable-by-Design Neural Feedback Control and the Youla-REN, arXiv (2025).

d-Tube: Why it's different



d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz



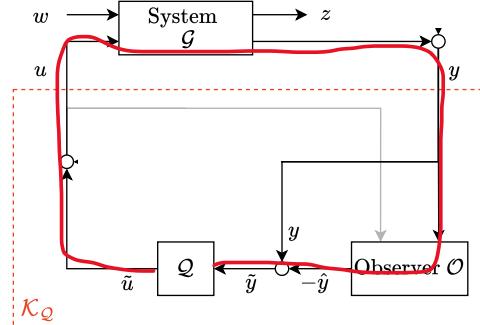
d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$



[20] Inspired by a conversation with Prof. A. Megretski one Friday afternoon in Boston (2023).

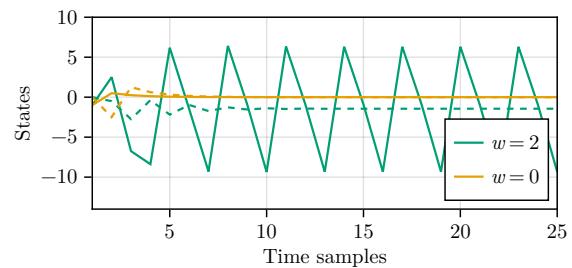
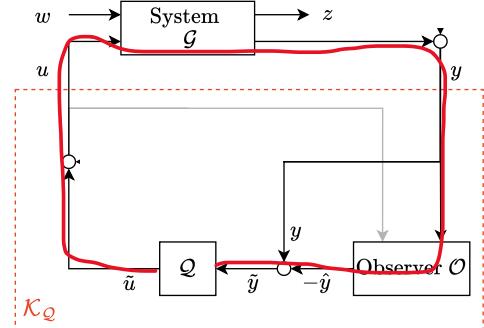
d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$



[20] Inspired by a conversation with Prof. A. Megretski one Friday afternoon in Boston (2023).

d-Tube: Why it's different

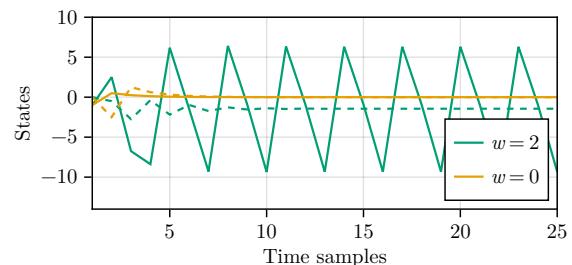
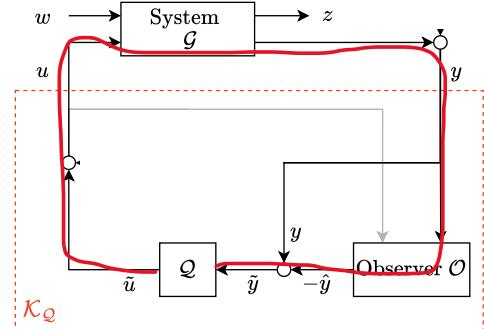
- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- What's happening?



[20] Inspired by a conversation with Prof. A. Megretski one Friday afternoon in Boston (2023).

d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

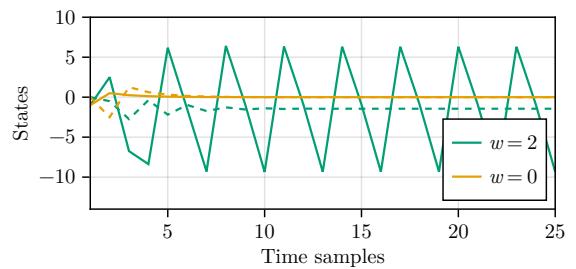
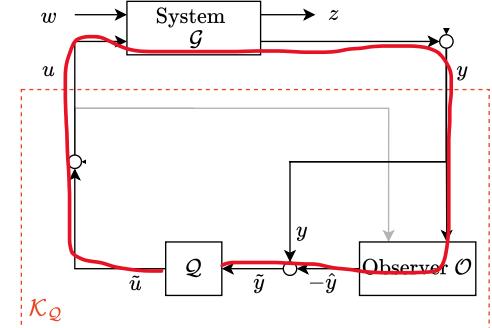
$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- What's happening?

$$\tilde{y}_t = y_t - \hat{y}_t = \tilde{x}_t$$

$$\tilde{x}_{t+1} = 0.5|\tilde{x}_t + \hat{x}_t| - 0.5|\hat{x}_t| + w_t$$



[20] Inspired by a conversation with Prof. A. Megretski one Friday afternoon in Boston (2023).

d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

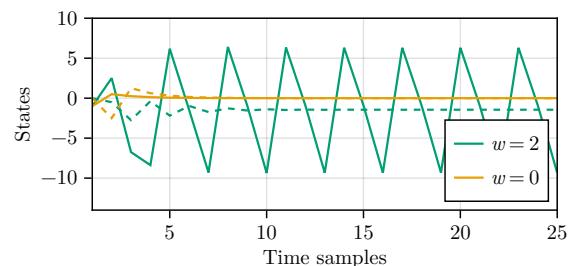
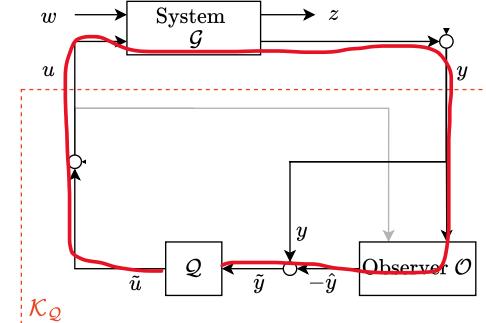
$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- What's happening?

$$\tilde{y}_t = y_t - \hat{y}_t = \tilde{x}_t$$

$$\tilde{x}_{t+1} = 0.5|\tilde{x}_t + \hat{x}_t| - 0.5|\hat{x}_t| + w_t$$

- \tilde{y} depends on u , which depends on \tilde{y} , etc...
- High-gain Q can de-stabilize the feedback loop



d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

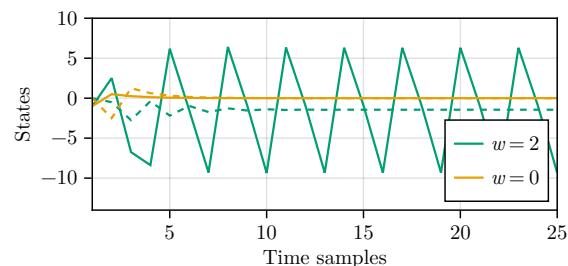
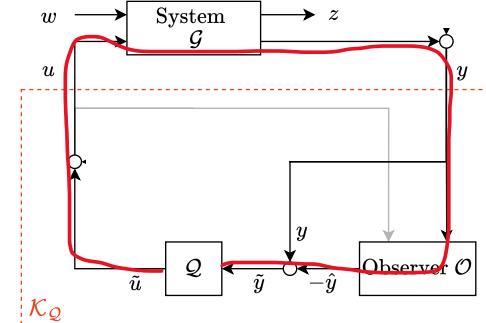
$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- What's happening?

$$\tilde{y}_t = y_t - \hat{y}_t = \tilde{x}_t$$

$$\tilde{x}_{t+1} = 0.5|\tilde{x}_t + \hat{x}_t| - 0.5|\hat{x}_t| + w_t$$

- \tilde{y} depends on u , which depends on \tilde{y} , etc...
- High-gain Q can de-stabilize the feedback loop



d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

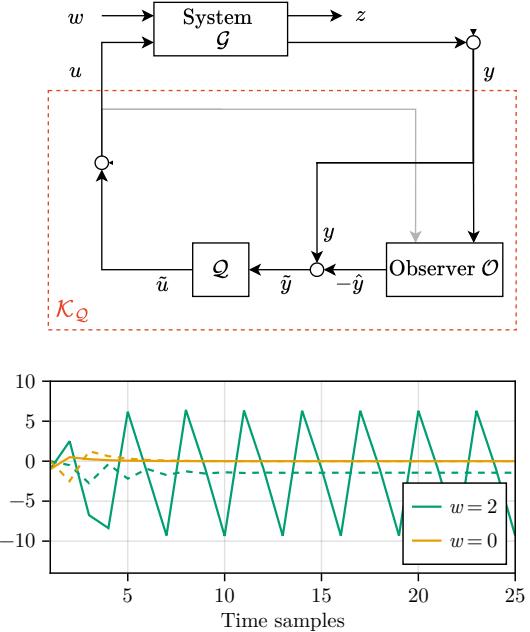
$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- What's happening?

$$\tilde{y}_t = y_t - \hat{y}_t = \tilde{x}_t$$

$$\tilde{x}_{t+1} = 0.5|\tilde{x}_t + \hat{x}_t| - 0.5|\hat{x}_t| + w_t$$

- **Linear case:** \hat{x} terms cancel out



d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

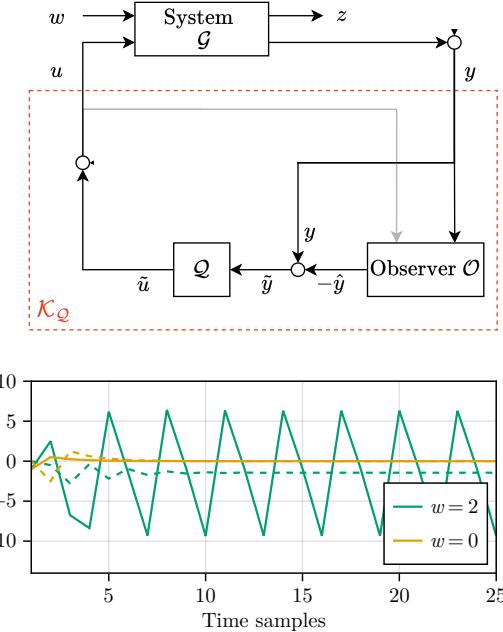
$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- What's happening?

$$\tilde{y}_t = y_t - \hat{y}_t = \tilde{x}_t$$

$$\tilde{x}_{t+1} = 0.5|\tilde{x}_t + \hat{x}_t| - 0.5|\hat{x}_t| + w_t$$

- **Linear case:** \hat{x} terms cancel out



d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5\underline{x}_t + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5\hat{x}_t + u_t, \quad \hat{y}_t = \hat{x}_t$$

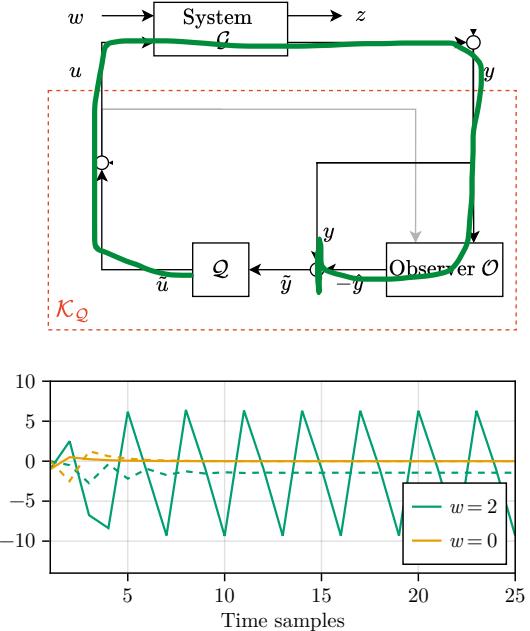
$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- What's happening?

$$\tilde{y}_t = y_t - \hat{y}_t = \tilde{x}_t$$

$$\tilde{x}_{t+1} = 0.5\tilde{x}_t + w_t$$

- **Linear case:** \hat{x} terms cancel out



d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5\underline{x}_t + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5\hat{x}_t + u_t, \quad \hat{y}_t = \hat{x}_t$$

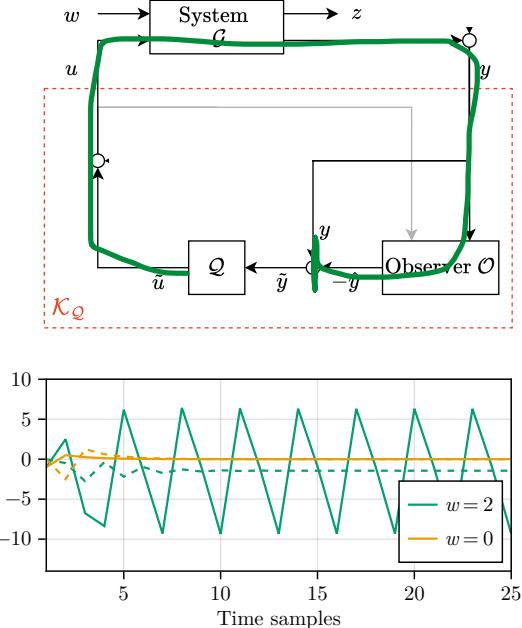
$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- What's happening?

$$\begin{aligned} \tilde{y}_t &= y_t - \hat{y}_t = \tilde{x}_t \\ \tilde{x}_{t+1} &= 0.5\tilde{x}_t + w_t \end{aligned}$$

Surprises have
no u dependence!

- Linear case: \hat{x} terms cancel out



d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5\underline{x}_t + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5\hat{x}_t + u_t, \quad \hat{y}_t = \hat{x}_t$$

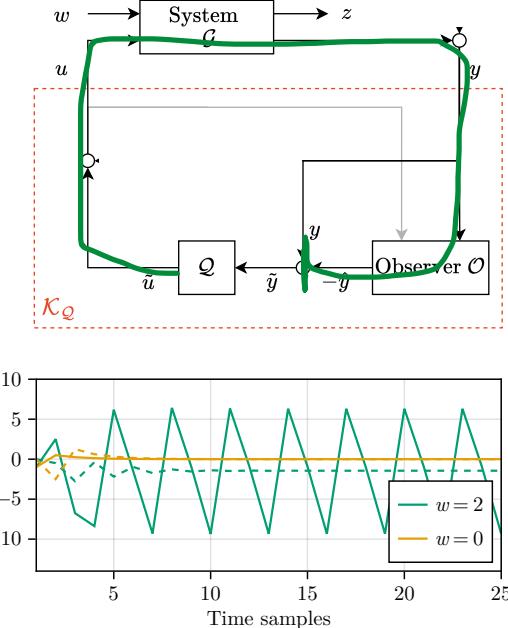
$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- What's happening?

$$\begin{aligned} \tilde{y}_t &= y_t - \hat{y}_t = \tilde{x}_t \\ \tilde{x}_{t+1} &= 0.5\tilde{x}_t + w_t \end{aligned}$$

Surprises have
no u dependence!

- Linear case: \hat{x} terms cancel out



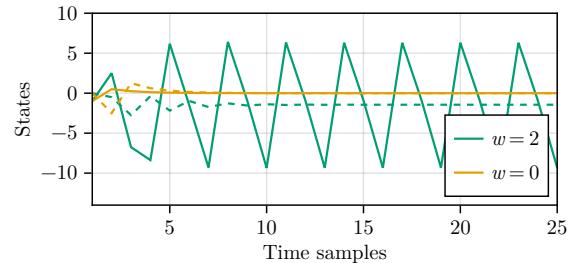
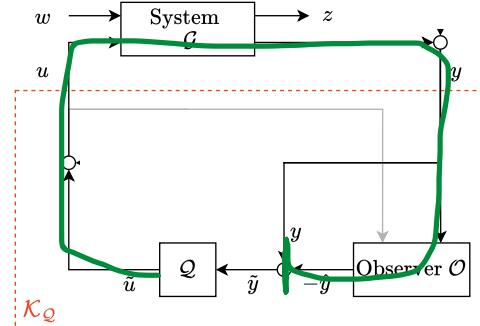
d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$



[20] Inspired by a conversation with Prof. A. Megretski one Friday afternoon in Boston (2023).

d-Tube: Why it's different

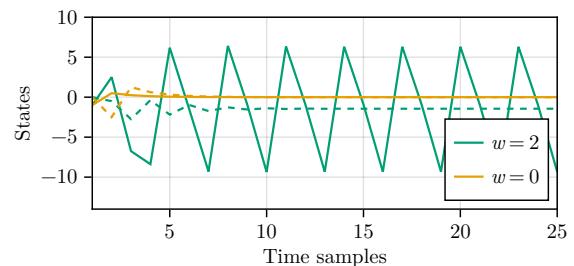
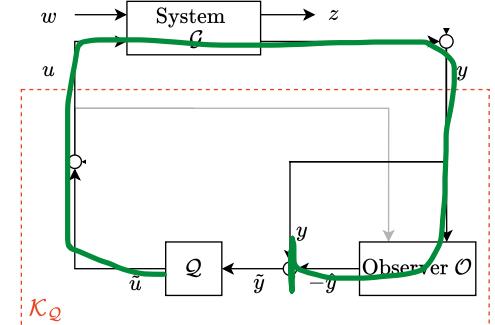
- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|x_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- Why's it still bounded for nonlinear?



[20] Inspired by a conversation with Prof. A. Megretski one Friday afternoon in Boston (2023).

d-Tube: Why it's different

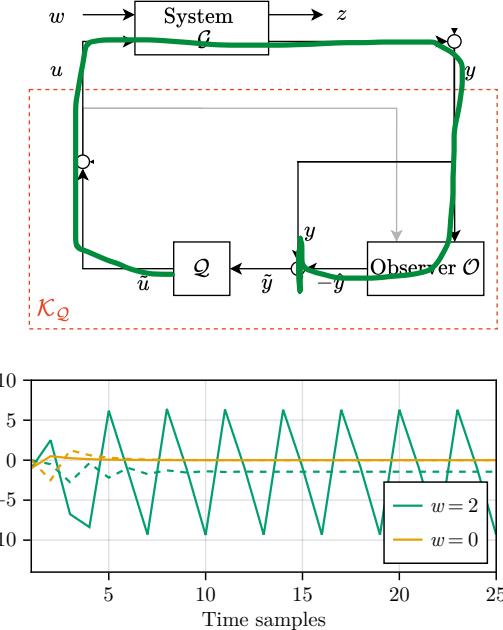
- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

$$\mathcal{G} : \quad x_{t+1} = 0.5|\underline{x}_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- Why's it still bounded for nonlinear?
- Can show: $\|\tilde{y}\|_T \leq \gamma_1 \|w\|_T + \gamma_2 |\tilde{x}_0|$



[20] Inspired by a conversation with Prof. A. Megretski one Friday afternoon in Boston (2023).

d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

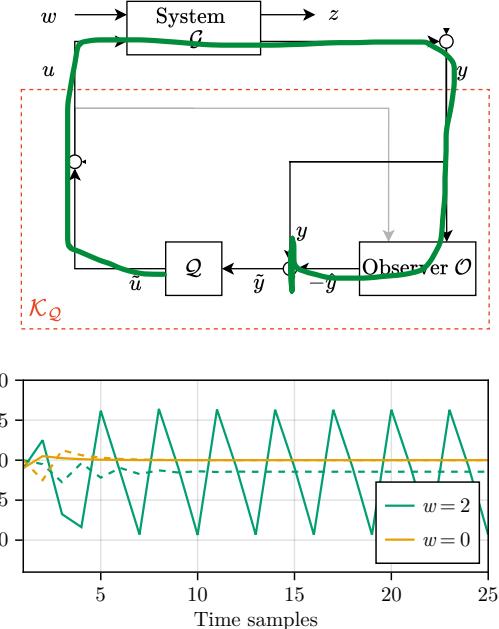
$$\mathcal{G} : \quad x_{t+1} = 0.5|\underline{x}_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- Why's it still bounded for nonlinear?
- Can show: $\|\tilde{y}\|_T \leq \gamma_1 \|w\|_T + \gamma_2 |\tilde{x}_0|$

Bound on surprises has
no u dependence!



d-Tube: Why it's different

- Linear systems: Youla \implies contracting & Lipschitz
- Nonlinear systems: Youla $\not\implies$ contracting & Lipschitz
- Example^[20]:

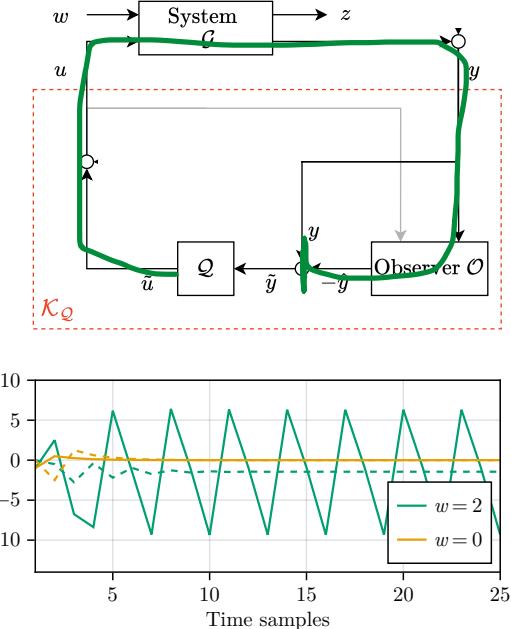
$$\mathcal{G} : \quad x_{t+1} = 0.5|\underline{x}_t| + u_t + w_t, \quad y_t = x_t$$

$$\mathcal{O} : \quad \hat{x}_{t+1} = 0.5|\hat{x}_t| + u_t, \quad \hat{y}_t = \hat{x}_t$$

$$u = \mathcal{Q}(\tilde{y}) := \min(2.5 - 5\tilde{y}, 0)$$

- Why's it still bounded for nonlinear?
- Can show: $\|\tilde{y}\|_T \leq \gamma_1 \|w\|_T + \gamma_2 |\tilde{x}_0|$

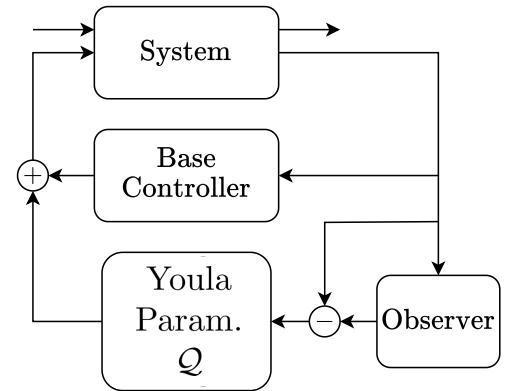
Bound on surprises has
no u dependence!



Outline

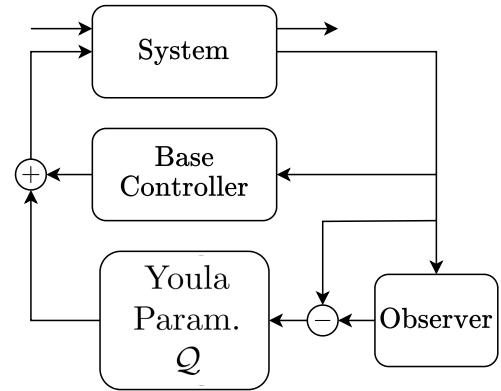
1. Motivation
2. Background: robust neural networks
3. Lipschitz-bounded policy networks
4. React to surprises with Youla-REN (Theory)
5. **React to surprises with Youla-REN (Application)**
6. Robust recurrent deep networks
7. Conclusions

The Youla-REN



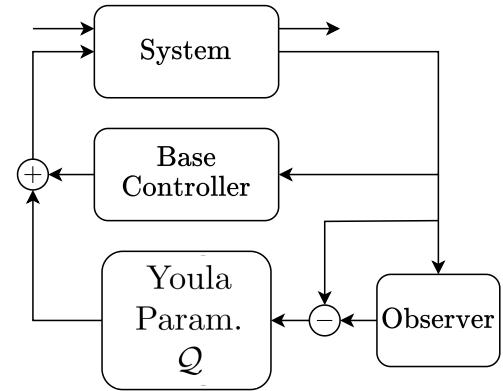
The Youla-REN

- Theory says: closed-loop stability guaranteed if we have **contracting & Lipschitz Q**



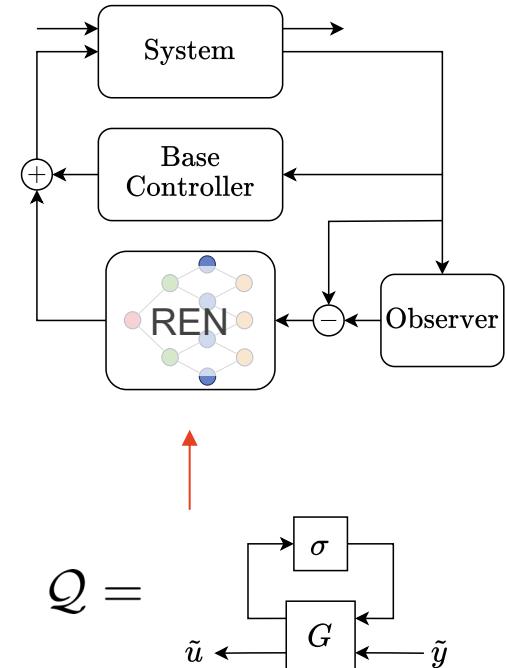
The Youla-REN

- Theory says: closed-loop stability guaranteed if we have **contracting & Lipschitz Q**
- Implementation: Recurrent Equilibrium Network (REN)



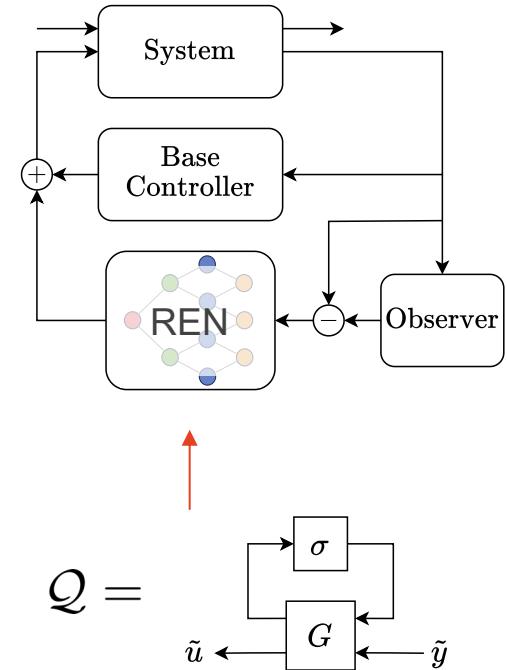
The Youla-REN

- Theory says: closed-loop stability guaranteed if we have **contracting & Lipschitz Q**
- Implementation: Recurrent Equilibrium Network (REN)



The Youla-REN

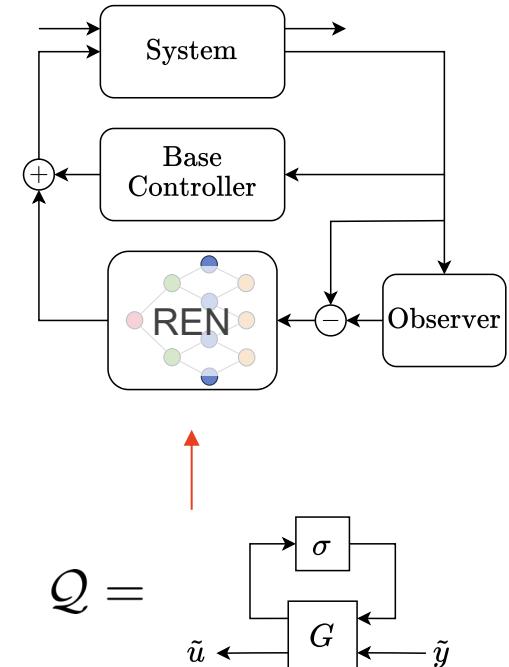
- Theory says: closed-loop stability guaranteed if we have **contracting & Lipschitz Q**
- Implementation: Recurrent Equilibrium Network (REN)
 - Contracting & Lipschitz
 - Universal approximator
 - **Direct parametrisation**



The Youla-REN

- Theory says: closed-loop stability guaranteed if we have **contracting & Lipschitz Q**
- Implementation: Recurrent Equilibrium Network (REN)
 - Contracting & Lipschitz
 - Universal approximator
 - **Direct parametrisation**

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools 😊

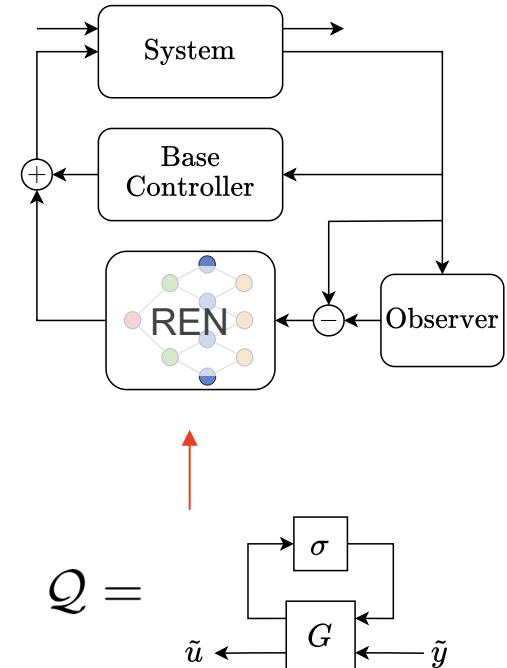


The Youla-REN

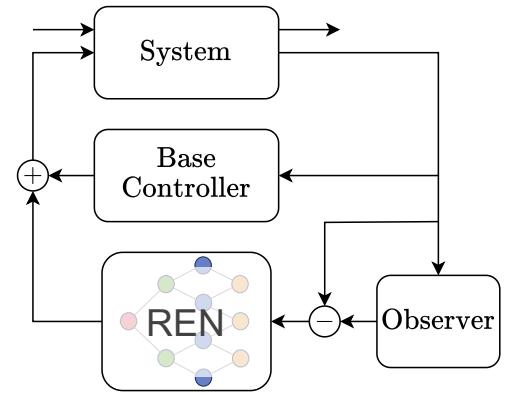
- Theory says: closed-loop stability guaranteed if we have **contracting & Lipschitz Q**
- Implementation: Recurrent Equilibrium Network (REN)
 - Contracting & Lipschitz
 - Universal approximator
 - **Direct parametrisation**

Directly learn controllers with robust stability guarantees using standard ML/deep RL tools 😊

- Practical implementation:
 - Design base controller and observer first
 - If model uncertainty, restrict gain of REN

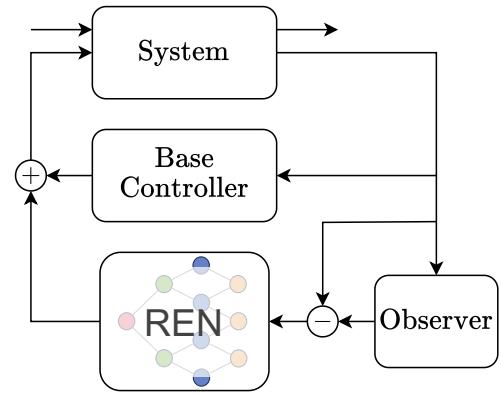


Learning with Stability: Example 1



Learning with Stability: Example 1

- Learning to stabilise nonlinear systems with economic costs



Learning with Stability: Example 1

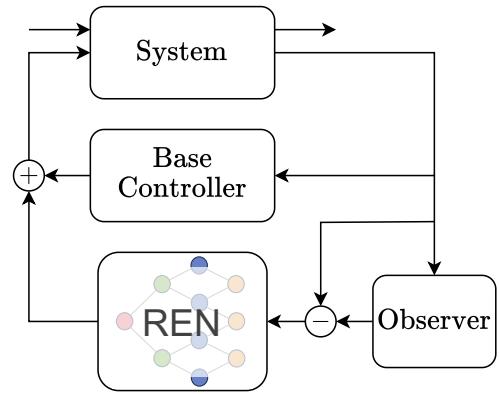
- Learning to stabilise nonlinear systems with economic costs
- System dynamics

$$\dot{x}_1 = -x_1 + x_3 + w_1$$

$$\dot{x}_2 = x_1^2 - x_2 - 2x_1x_3 + x_3 + w_2$$

$$\dot{x}_3 = -x_2 + u + w_3$$

$$y = [x_2; x_3] + v$$



Learning with Stability: Example 1

- Learning to stabilise nonlinear systems with economic costs
- System dynamics

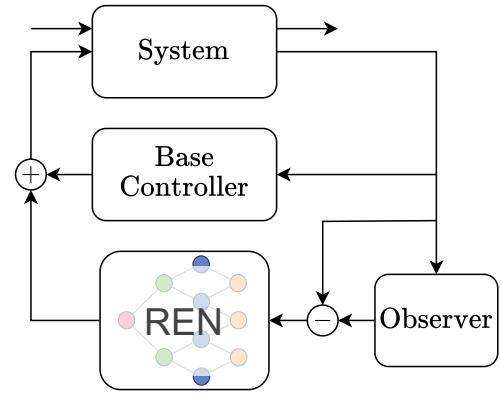
$$\dot{x}_1 = -x_1 + x_3 + w_1$$

$$\dot{x}_2 = x_1^2 - x_2 - 2x_1x_3 + x_3 + w_2$$

$$\dot{x}_3 = -x_2 + u + w_3$$

$$y = [x_2; x_3] + v$$

- Cost function
- $$\min \sum_t (|u_t| + 500 \max(|u_t| - u_{\max}, 0))$$
- Optimal policy is $u = 0$ (closed-loop unstable)



Learning with Stability: Example 1

- Learning to stabilise nonlinear systems with economic costs
- System dynamics

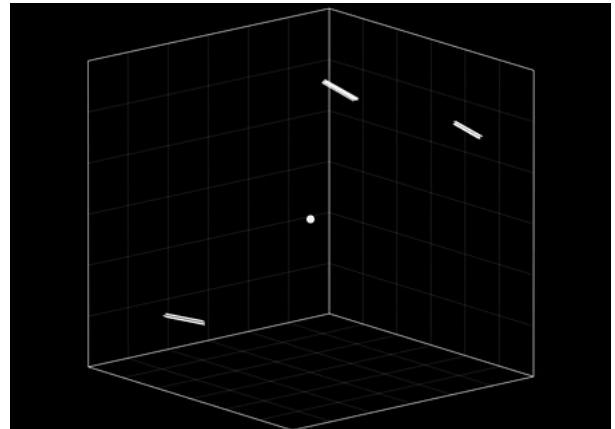
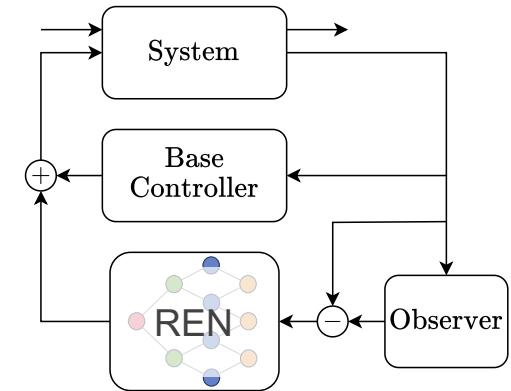
$$\dot{x}_1 = -x_1 + x_3 + w_1$$

$$\dot{x}_2 = x_1^2 - x_2 - 2x_1x_3 + x_3 + w_2$$

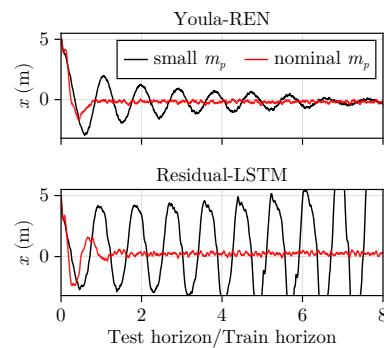
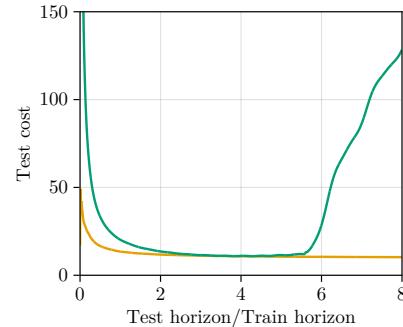
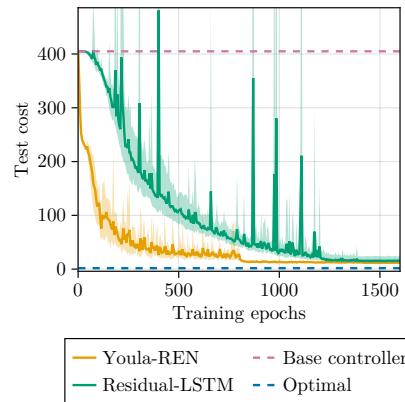
$$\dot{x}_3 = -x_2 + u + w_3$$

$$y = [x_2; x_3] + v$$

- Cost function
- $$\min \sum_t (|u_t| + 500 \max(|u_t| - u_{\max}, 0))$$
- Optimal policy is $u = 0$ (closed-loop unstable)

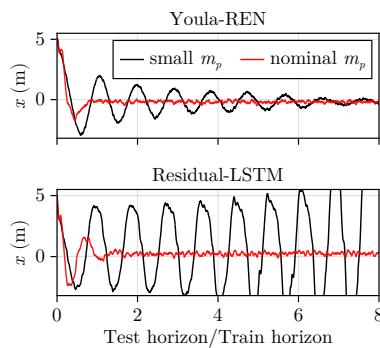
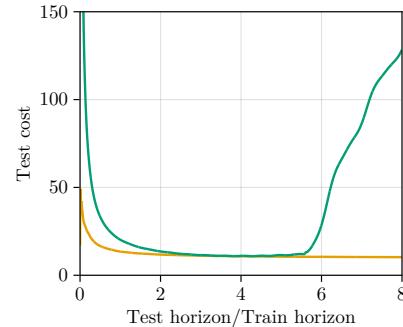
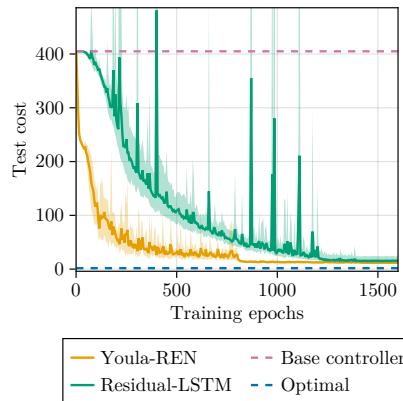


Learning with Stability: Example 2



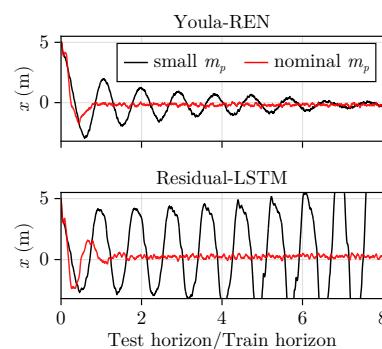
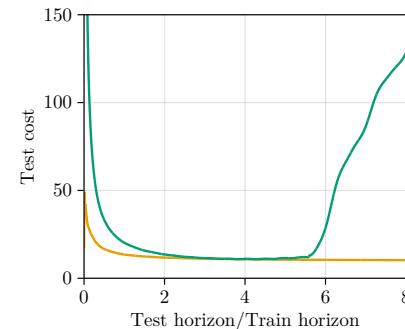
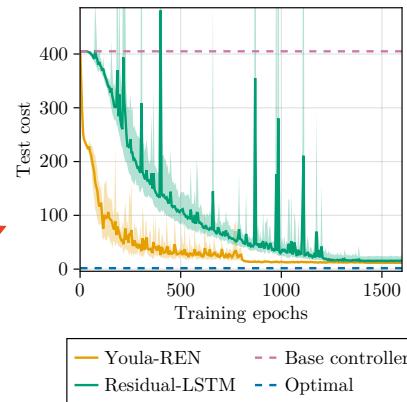
Learning with Stability: Example 2

- Linearised cart-pole, uncertain pole mass



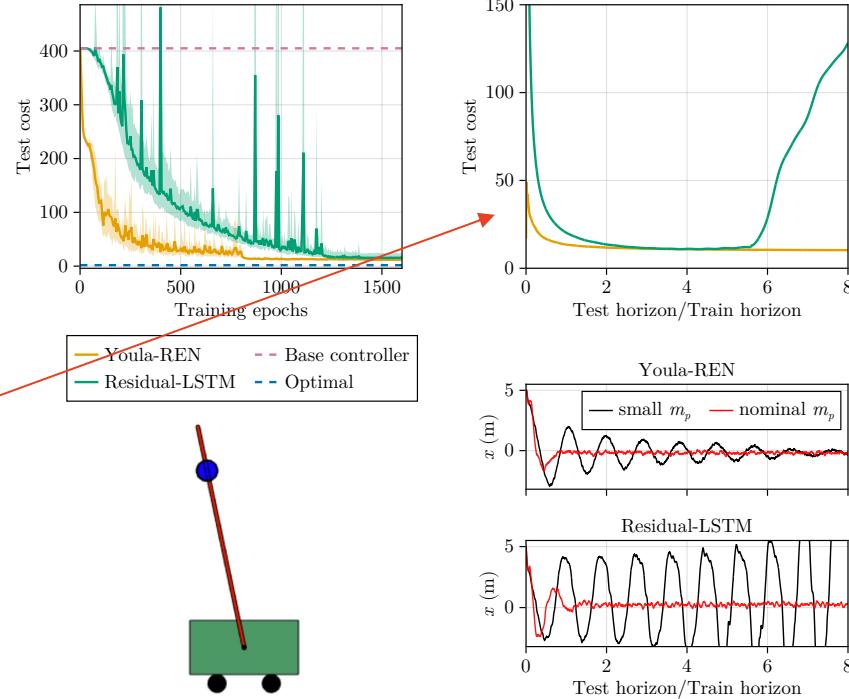
Learning with Stability: Example 2

- Linearised cart-pole, uncertain pole mass
- Youla-REN is closed-loop stable during training



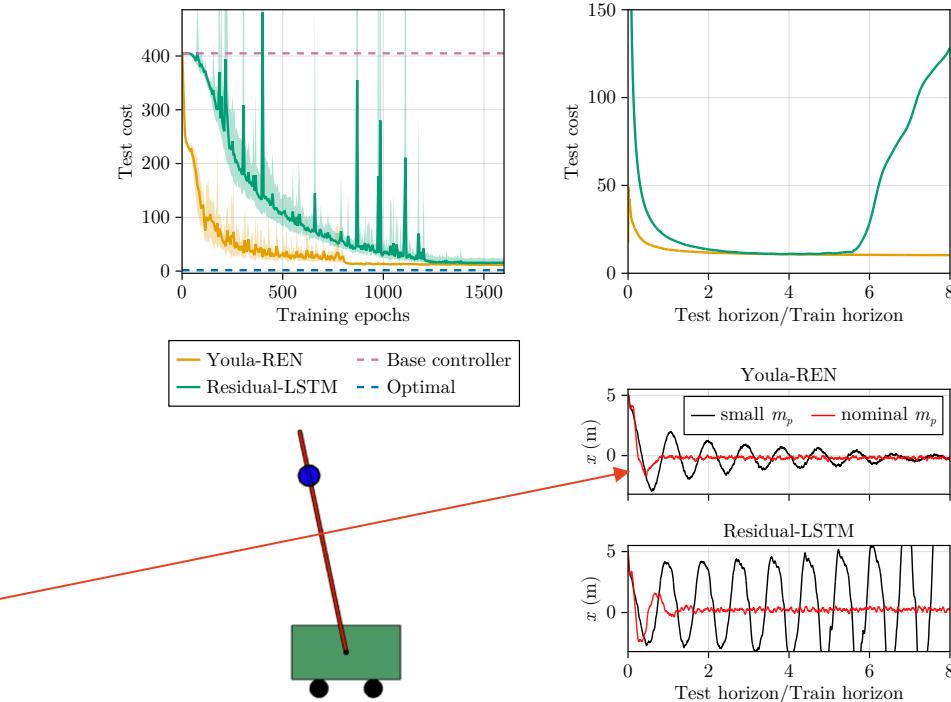
Learning with Stability: Example 2

- Linearised cart-pole, uncertain pole mass
- Youla-REN is closed-loop stable during training
- Youla-REN is closed-loop stable for long test horizons



Learning with Stability: Example 2

- Linearised cart-pole, uncertain pole mass
- Youla-REN is closed-loop stable during training
- Youla-REN is closed-loop stable for long test horizons
- Youla-REN “adapts” to the uncertain parameter



Outline

1. Motivation
2. Background: robust neural networks
3. Lipschitz-bounded policy networks
4. React to surprises with Youla-REN (Theory)
5. React to surprises with Youla-REN (Application)
6. Robust recurrent deep networks
7. Conclusions

Why Robust Recurrent Deep Network (R2DN)?

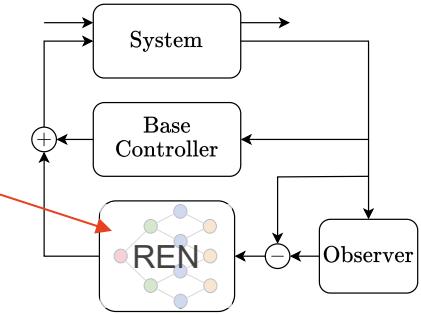
[21] Barbara et al., R2DN: Scalable Parametrization of Contracting and Lipschitz Recurrent Deep Networks, *arXiv* (2025).

The University of Sydney

35

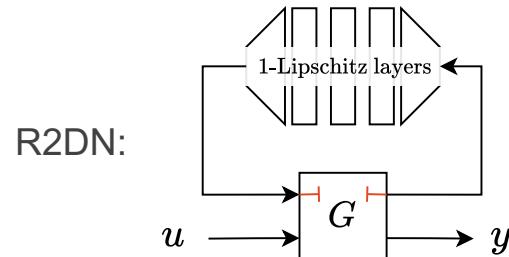
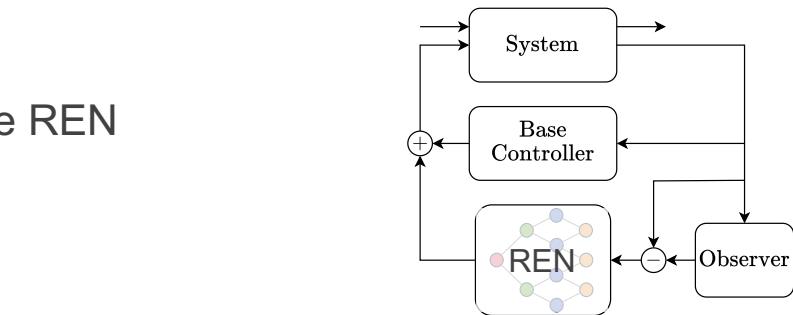
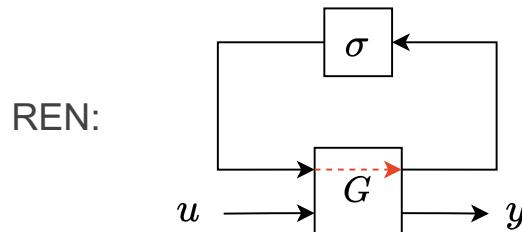
Why Robust Recurrent Deep Network (R2DN)?

- Computational bottleneck in Youla-REN is the REN



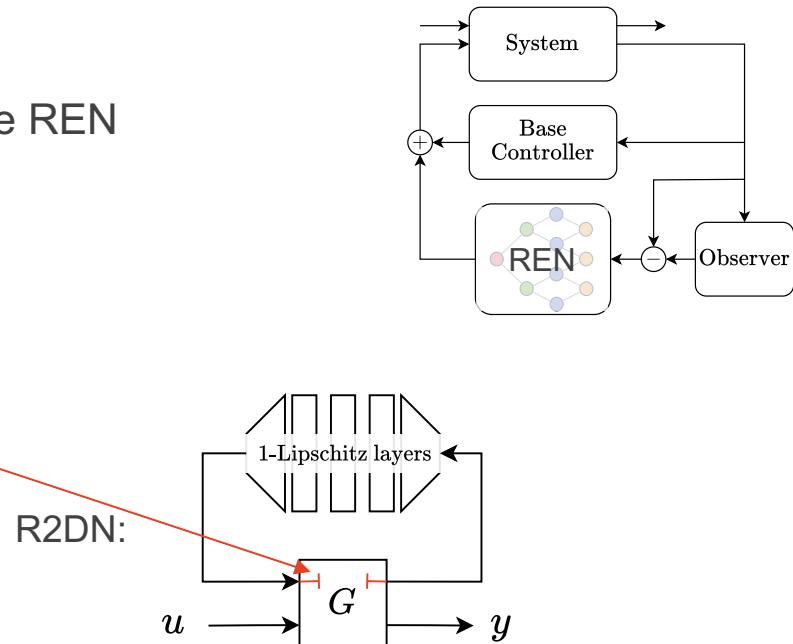
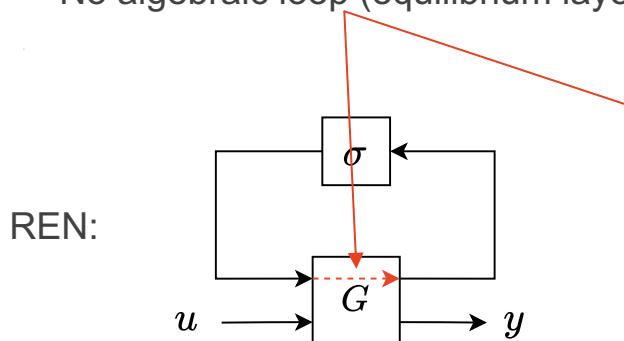
Why Robust Recurrent Deep Network (R2DN)?

- Computational bottleneck in Youla-REN is the REN
- Introducing R2DN – like REN, but:



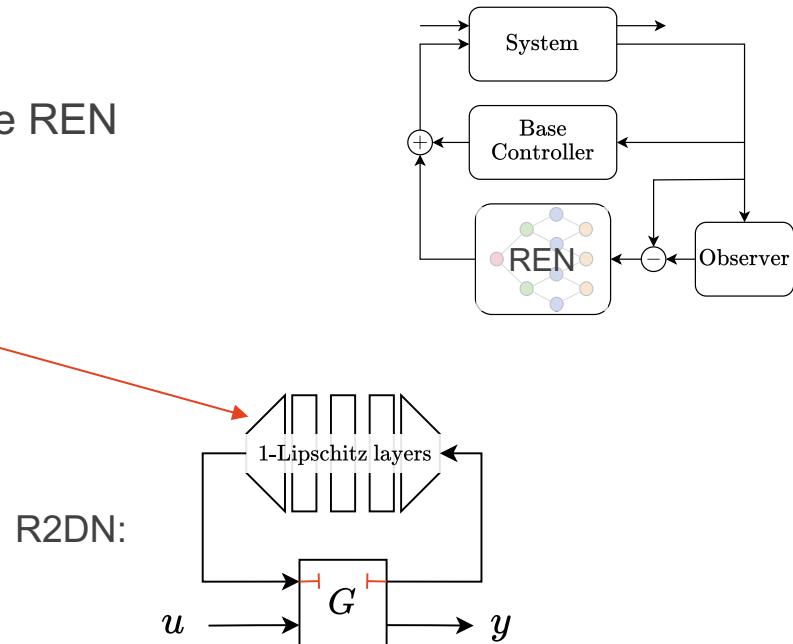
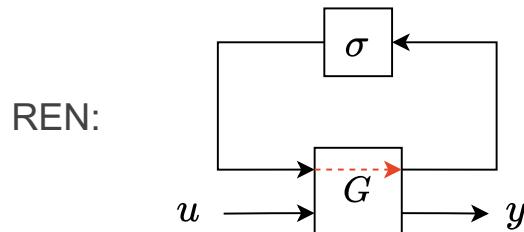
Why Robust Recurrent Deep Network (R2DN)?

- Computational bottleneck in Youla-REN is the REN
- Introducing R2DN – like REN, but:
 - No algebraic loop (equilibrium layer)



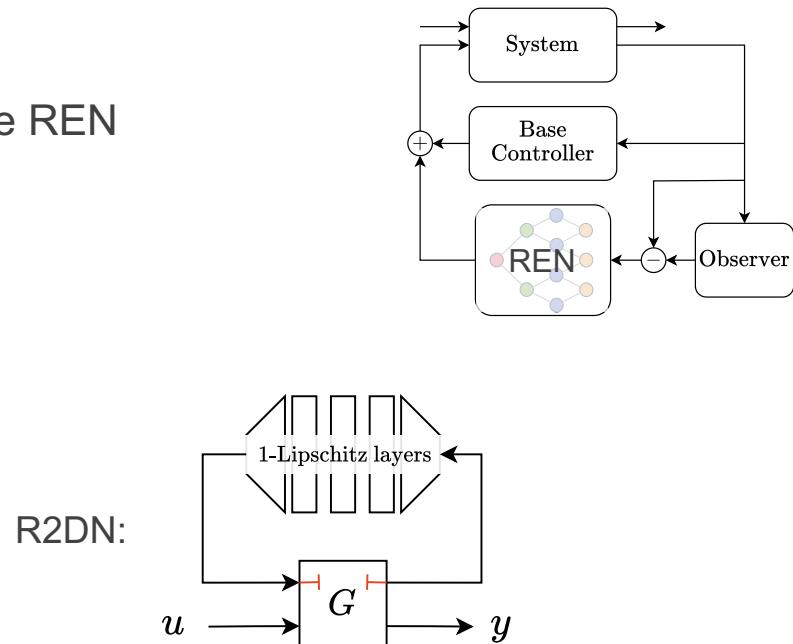
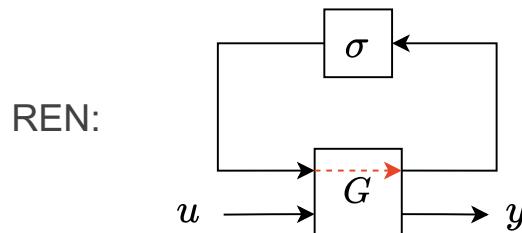
Why Robust Recurrent Deep Network (R2DN)?

- Computational bottleneck in Youla-REN is the REN
- Introducing R2DN – like REN, but:
 - No algebraic loop (equilibrium layer)
 - Nonlinearity can be any 1-Lipschitz network



Why Robust Recurrent Deep Network (R2DN)?

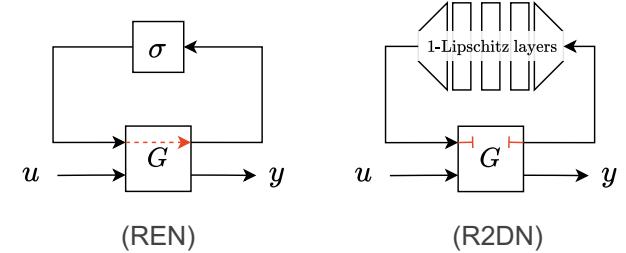
- Computational bottleneck in Youla-REN is the REN
- Introducing R2DN – like REN, but:
 - No algebraic loop (equilibrium layer)
 - Nonlinearity can be any 1-Lipschitz network



- Main contribution^[21]: direct parametrisation of contracting & Lipschitz R2DNs

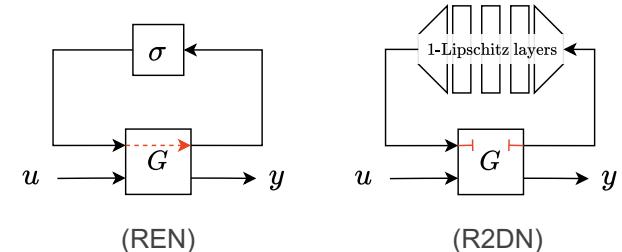
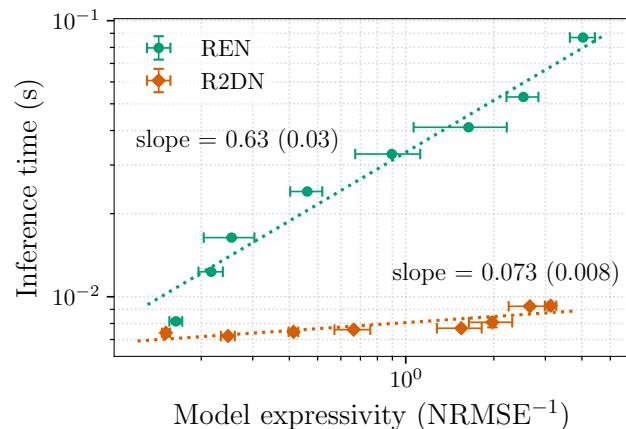
[21] Barbara et al., R2DN: Scalable Parametrization of Contracting and Lipschitz Recurrent Deep Networks, arXiv (2025).

Better Scalability, Similar Performance



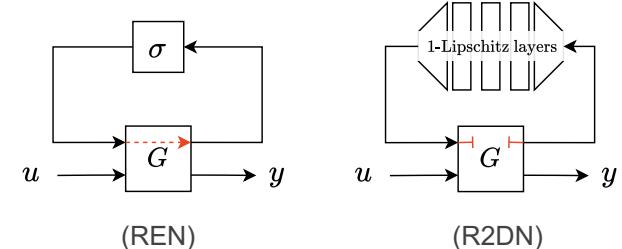
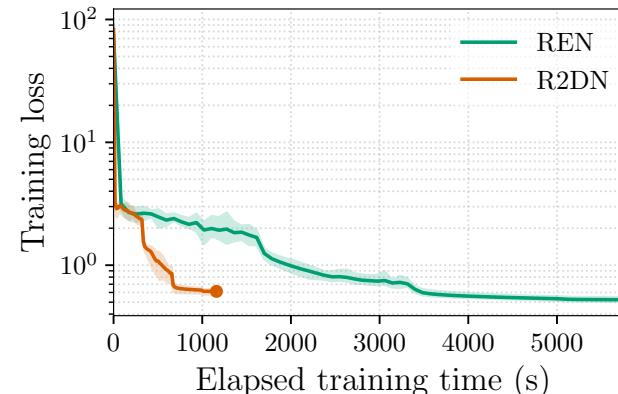
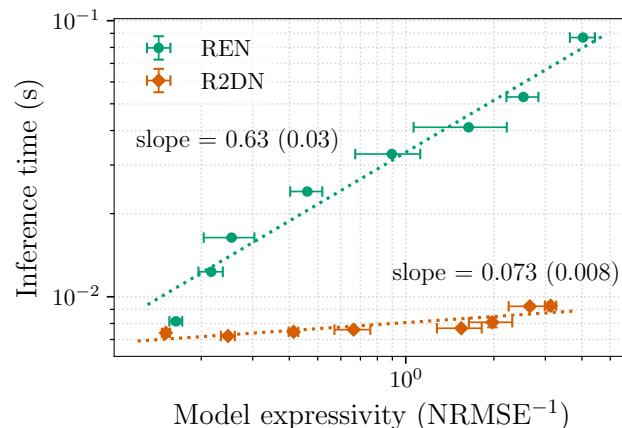
Better Scalability, Similar Performance

- Main benefits:
 - No algebraic loop → **efficient** GPU computation
 - No algebraic loop → **scalable**
 - Any 1-Lipschitz network → **flexible**



Better Scalability, Similar Performance

- Main benefits:
 - No algebraic loop → **efficient** GPU computation
 - No algebraic loop → **scalable**
 - Any 1-Lipschitz network → **flexible**
- Possible limitation (in theory): more conservative. In practice...



Outline

1. Motivation
2. Background: robust neural networks
3. Lipschitz-bounded policy networks
4. React to surprises with Youla-REN (Theory)
5. React to surprises with Youla-REN (Application)
6. Robust recurrent deep networks
7. Conclusions

Conclusions

- Robust NNs → learning-based control with stability & robustness **guarantees**
- The right choice of policy **parametrisation** and **architecture** enables:
 1. Closed-loop stability certificates
 2. Smooth, unconstrained parametrisations
 3. Non-restrictive parametrisations (under assumptions)
- “Plug-and-play” with standard deep RL and machine-learning pipelines

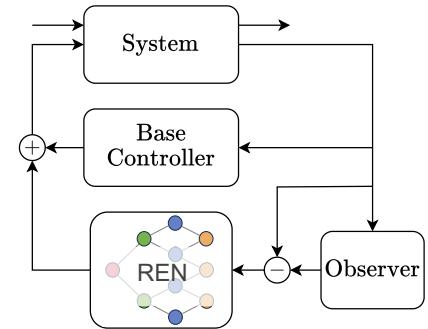
Future Work

Future Work

- Robust neural networks:
 - Develop more scalable parametrisations
 - Bigger library of **behavioural properties** & **network architectures**

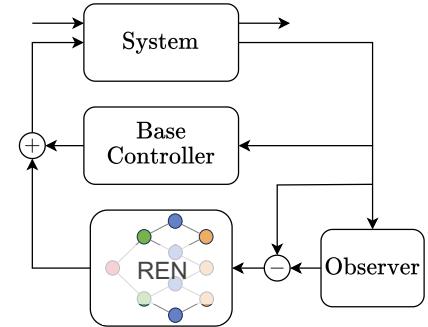
Future Work

- Robust neural networks:
 - Develop more scalable parametrisations
 - Bigger library of **behavioural properties & network architectures**
- Youla parametrisation:
 - Local stability
 - Non-smooth systems
 - Prescribed robustness
 - Designing good observers?



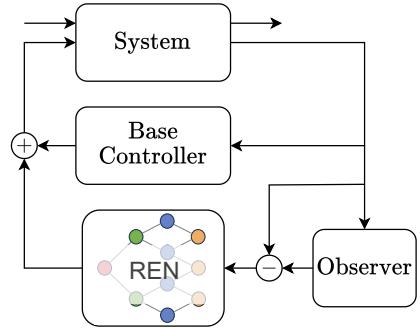
Future Work

- Robust neural networks:
 - Develop more scalable parametrisations
 - Bigger library of **behavioural properties** & **network architectures**
- Youla parametrisation:
 - Local stability
 - Non-smooth systems
 - Prescribed robustness
 - Designing good observers?
- Interplay between **architectures** and their **parametrisations**



Future Work

- Robust neural networks:
 - Develop more scalable parametrisations
 - Bigger library of **behavioural properties** & **network architectures**
- Youla parametrisation:
 - Local stability
 - Non-smooth systems
 - Prescribed robustness
 - Designing good observers?
- Interplay between **architectures** and their **parametrisations**
- Applications to robotics^[22]



[22] Pries & Ryll, Learning Robust Agile Flight Control with Stability Guarantees, arXiv (2025).